

Graph Algorithms

Weighted Graphs

Weighted Graphs

- Recall the graph

Weighted Graphs

- Recall the graph
 - $G = (V, E)$
 - V : Set of vertices
 - E : Set of edges
 - E is a subset of pairs (v, v') : $E \subseteq V \times V$
 - Undirected graph: (v, v') and (v', v) are same edge
 - Directed graph:
 - (v, v') is an edge from v to v'
 - Does not guarantee that (v', v) is also an edge

Weighted Graphs ...

- Adding edge weights

Weighted Graphs ...

- Adding edge weights
 - Label each edge with a number - **cost**

Weighted Graphs ...

- Adding edge weights
 - Label each edge with a number – **cost**
 - Ticket price on a flight sector
 - Tolls on highway segment
 - Distance travelled between two stations
 - Typical time between two location during peak hour traffic

Weighted Graphs ...

- Weighted graph
 - $G = (V, E)$
 - **Weight function**, $w: E \rightarrow \text{Reals}$

Weighted Graphs ...(Shortest paths)

- **Weighted graph**
 - $G = (V, E)$
 - **Weight function**, $w: E \rightarrow \text{Reals}$
- Let $e_1 = (v_0, v_1), e_2 = (v_1, v_2), \dots, e_n = (v_{n-1}, v_n)$ be a path from v_0 to v_n
- Cost of the path is $w(e_1) + w(e_2) + \dots + w(e_n)$
- **Shortest path** from v_0 to v_n : minimum cost

Shortest path problems

Shortest path problems

- **Single source**
 - Find shortest paths from some fixed vertex, say 1, to every other vertex

Shortest path problems

- **Single source**
 - Find shortest paths from some fixed vertex, say 1, to every other vertex
 - Transport finished product from factory (single source) to all retail outlets

Shortest path problems

- **Single source**

- Find shortest paths from some fixed vertex, say 1, to every other vertex
 - Transport finished product from factory (single source) to all retail outlets
 - Courier company delivers items from distribution center (single source) to addresses

Shortest path problems

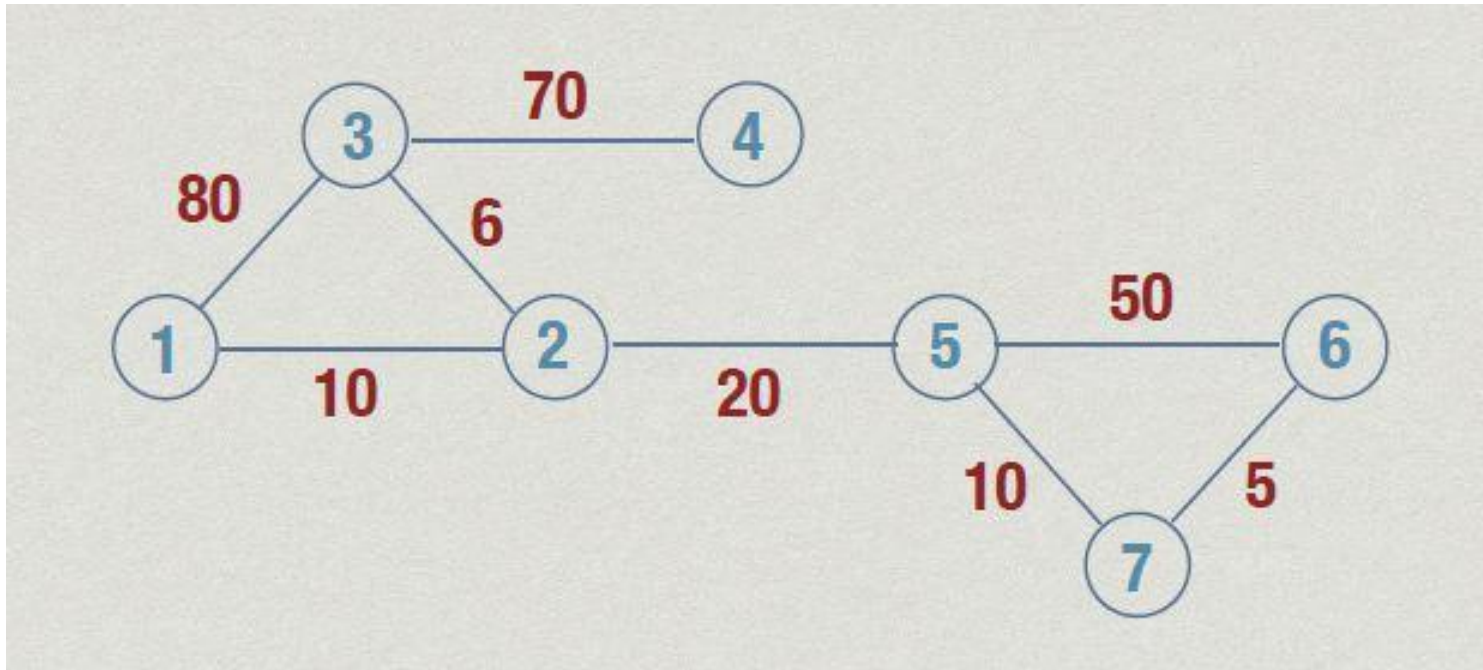
- **All pairs**
 - Find shortest paths between every pair of vertices i and j

Shortest path problems

- **All pairs**
 - Find shortest paths between every pair of vertices i and j
 - Railway routes, shortest way to travel between any pair of cities

Single source shortest paths

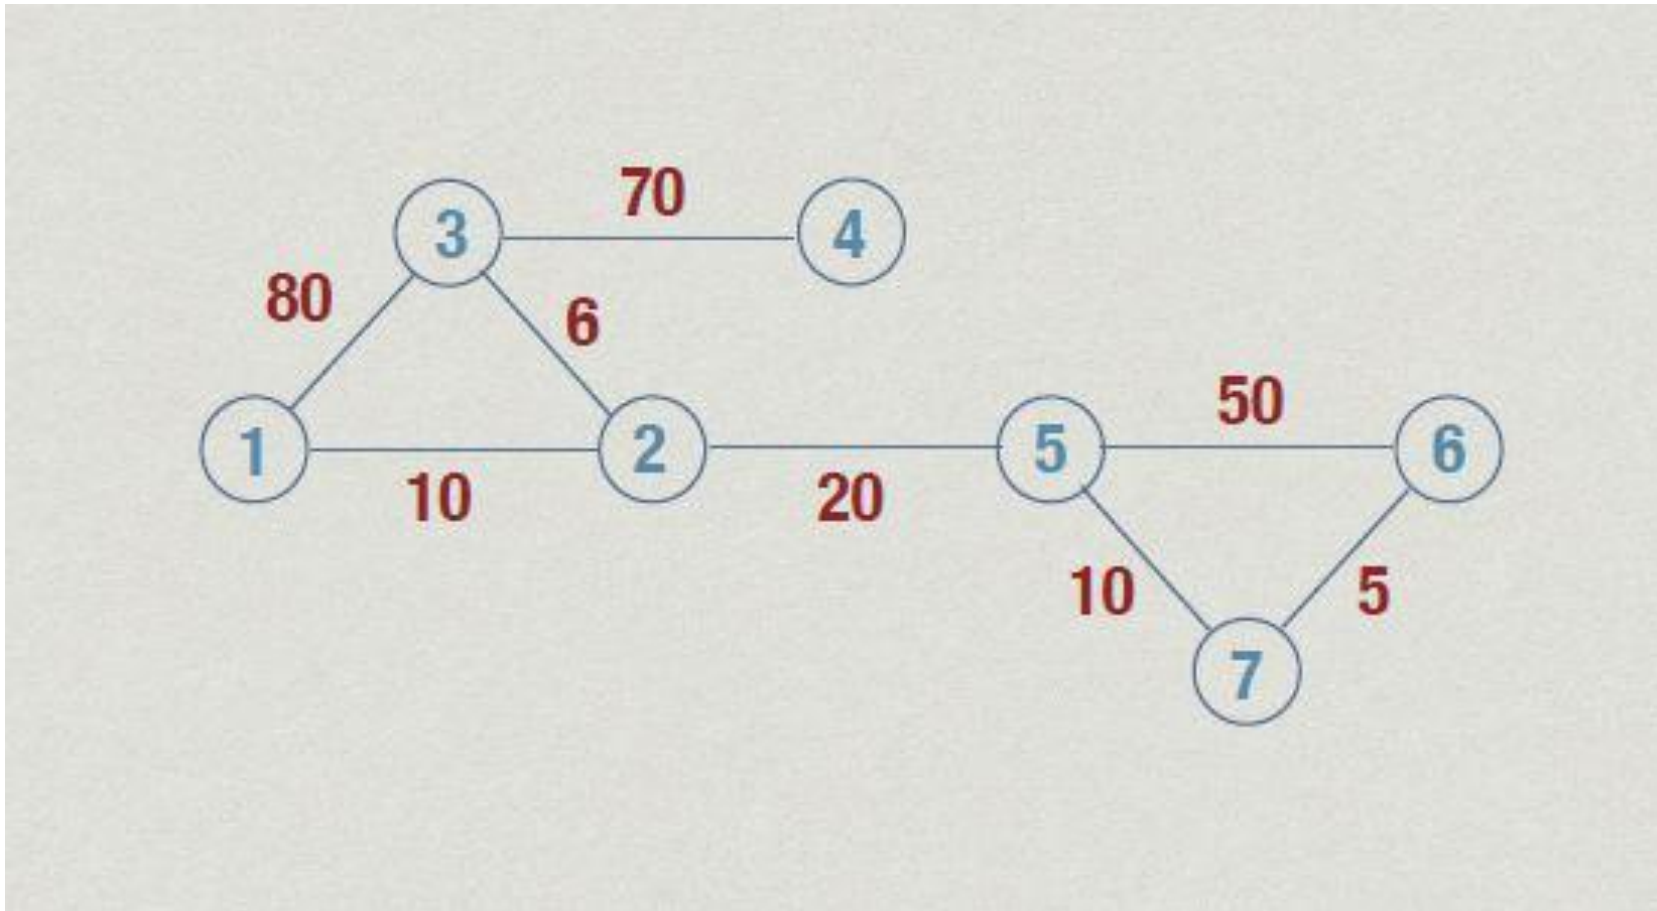
- For instance, shortest path from 1 to 2,3, ..., 7



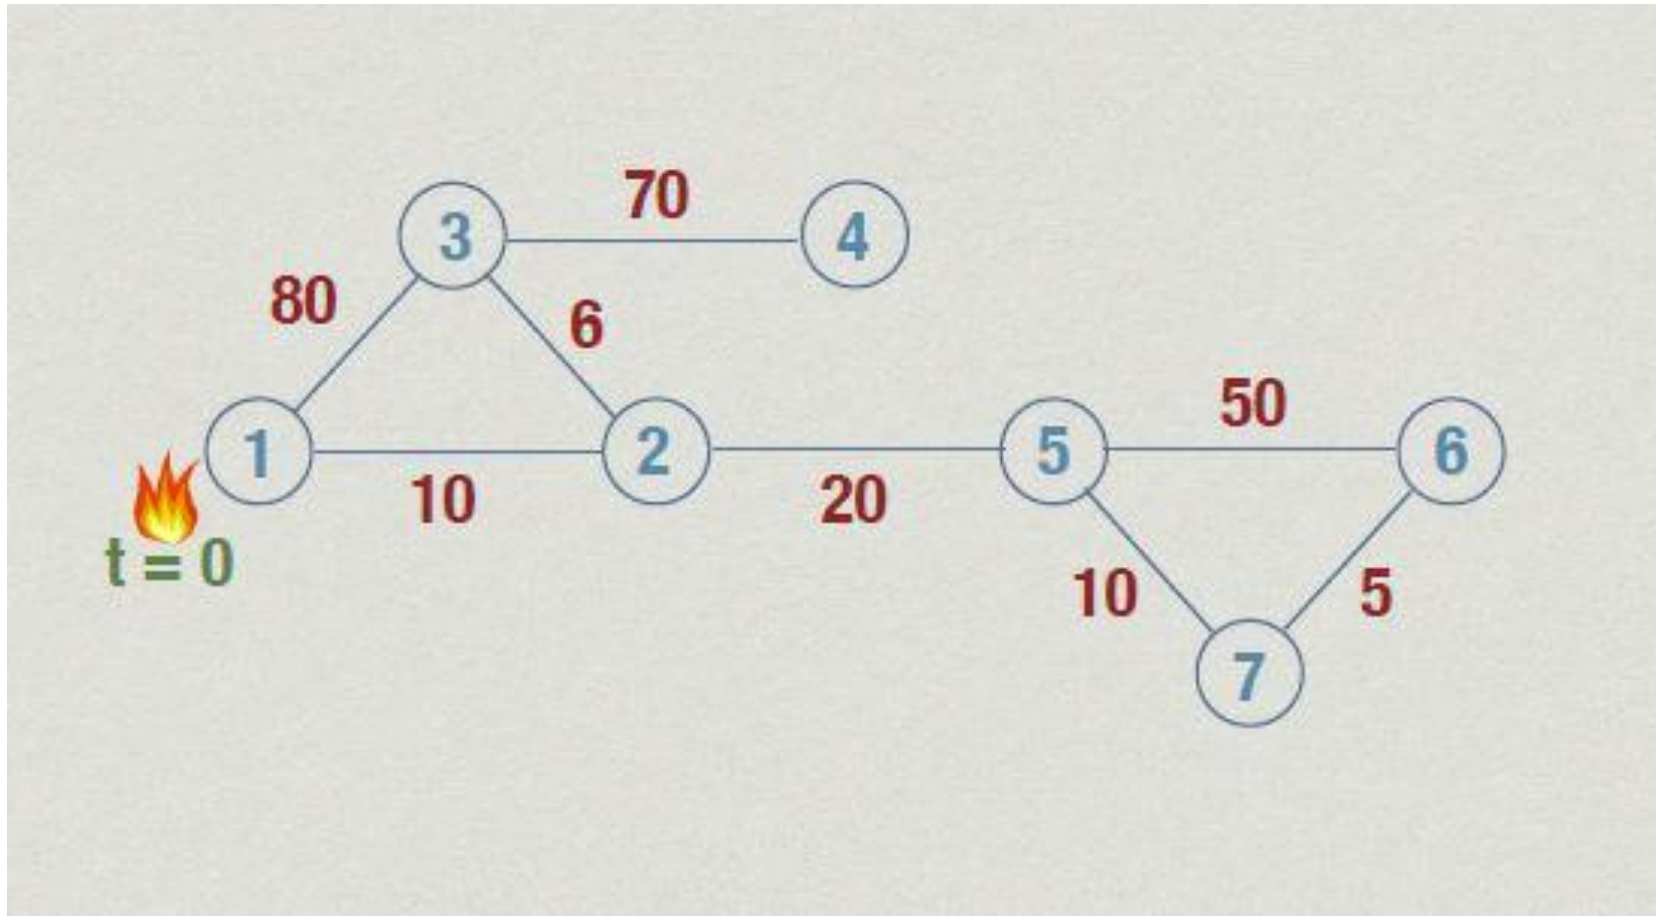
Single source shortest paths ...

- Imagine vertices are oil depots, edges are pipelines
- Set fire to oil depot at vertex 1
 - Fire travels at uniform speed along each pipeline
- First oil depot to catch fire after 1 is nearest vertex
- Next oil depot is second nearest vertex
- ...

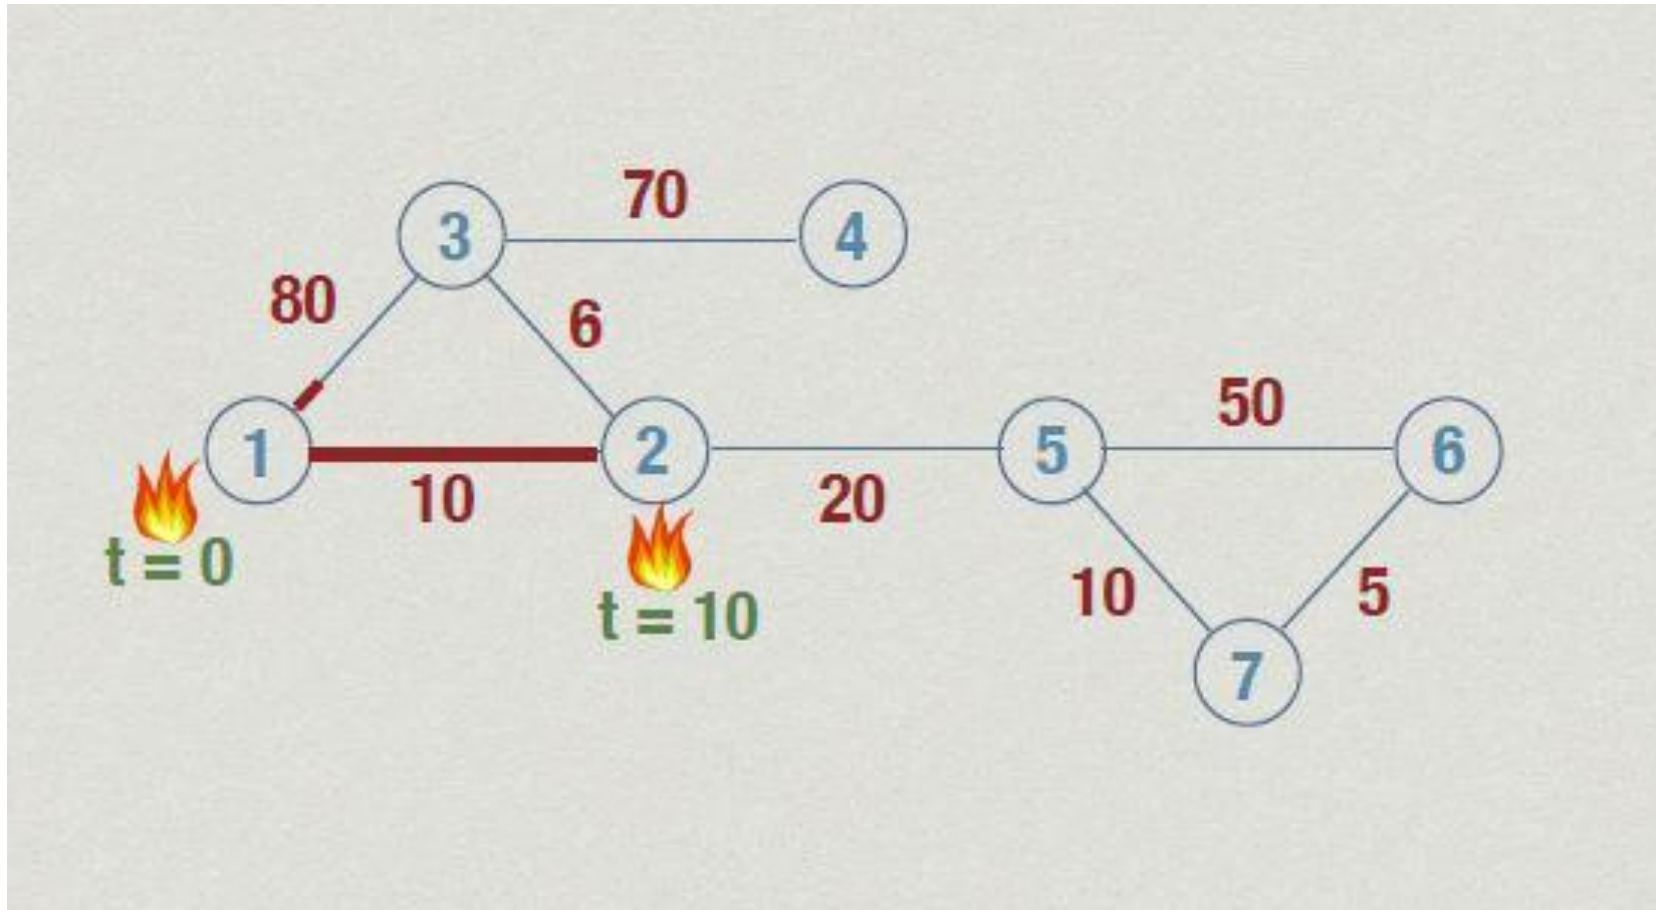
Single source shortest paths ...



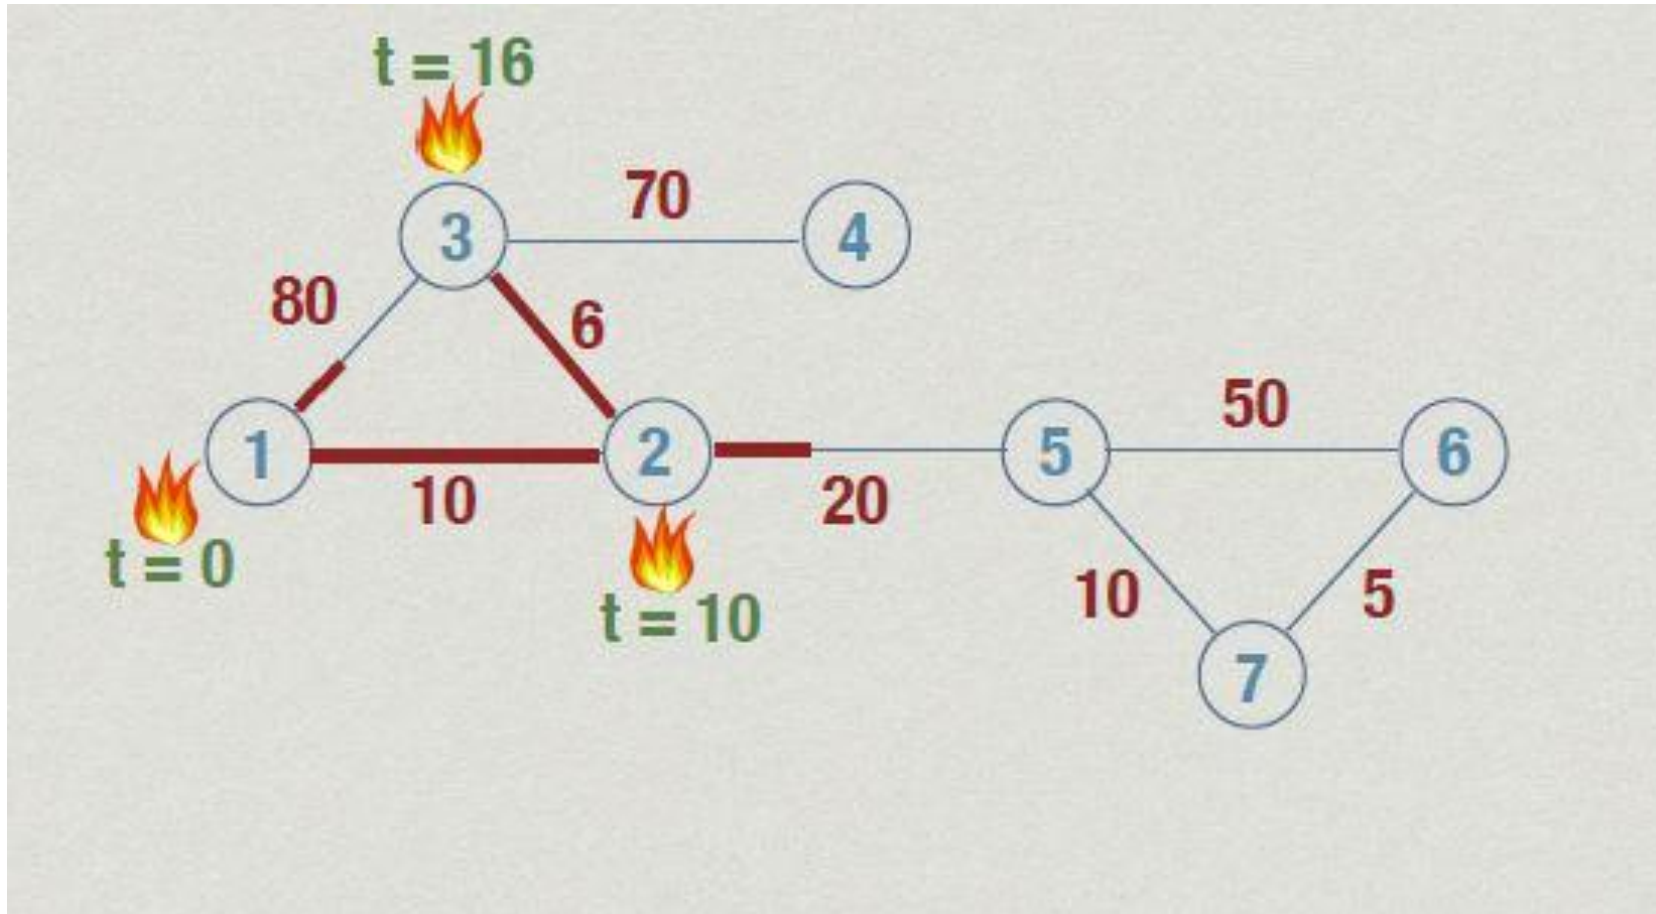
Single source shortest paths ...



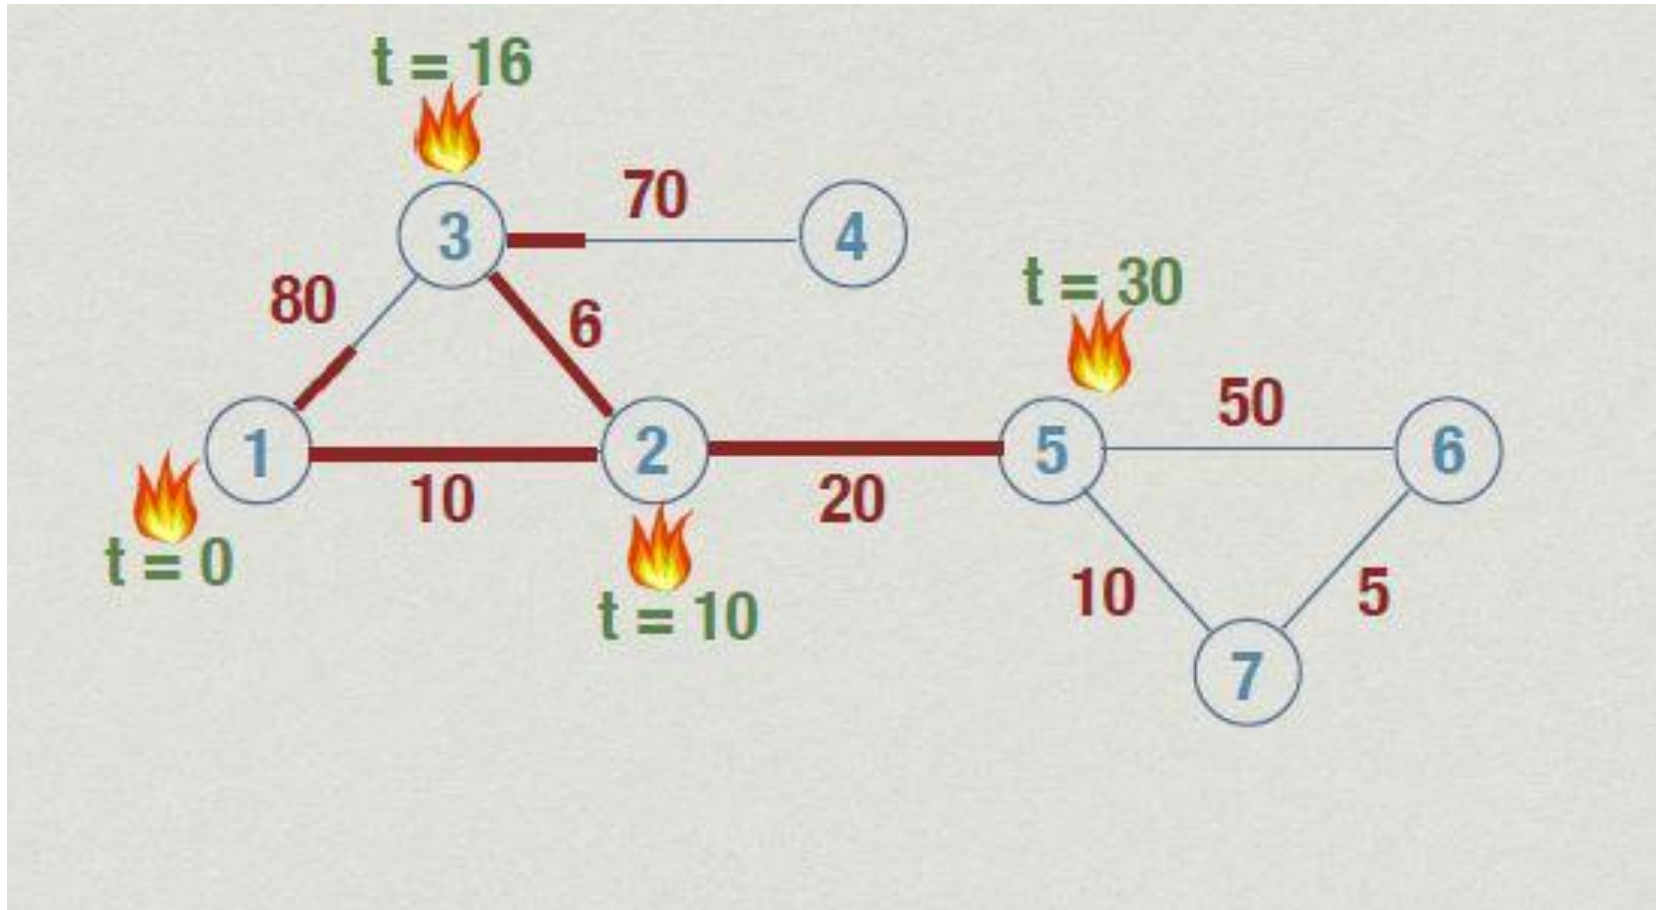
Single source shortest paths ...



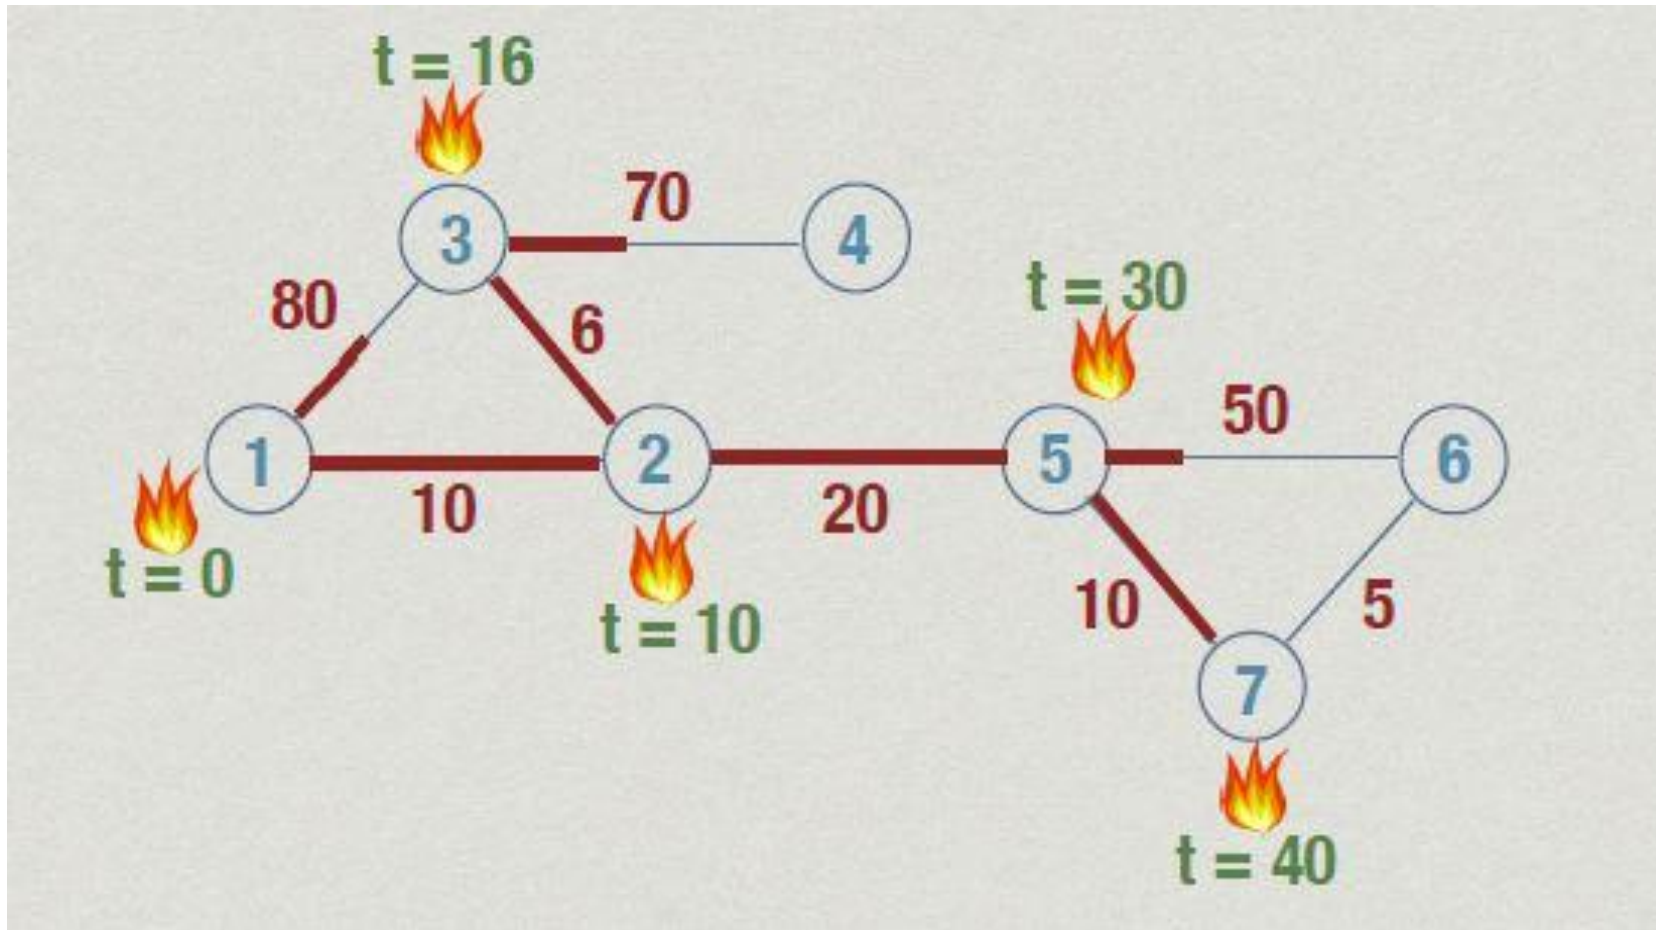
Single source shortest paths ...



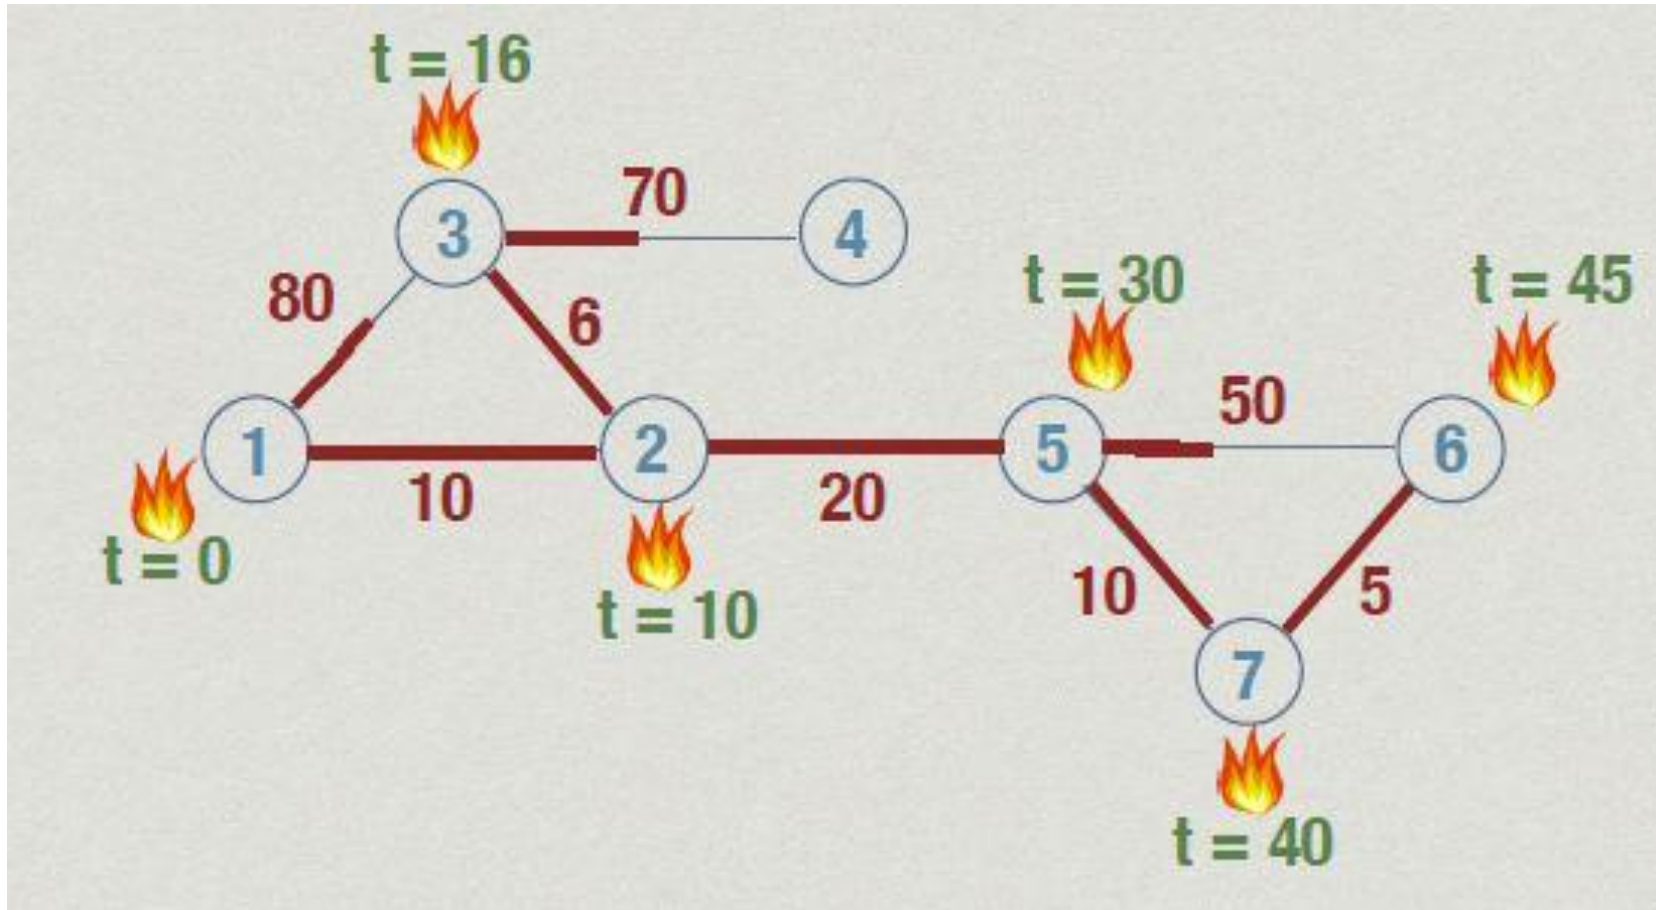
Single source shortest paths ...



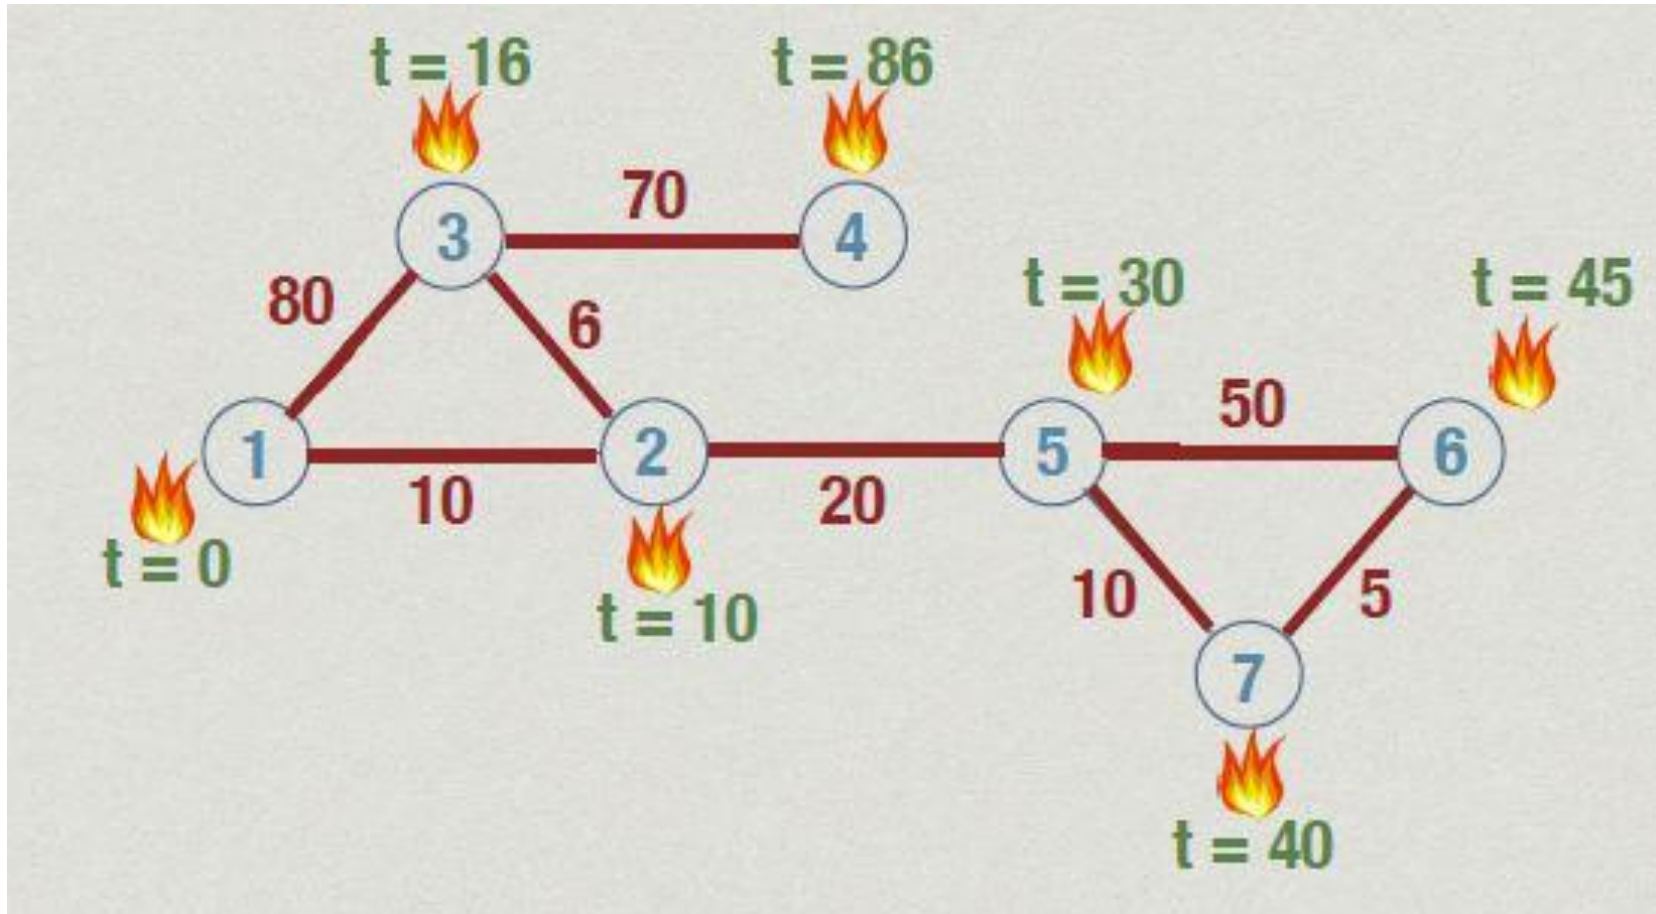
Single source shortest paths ...



Single source shortest paths ...



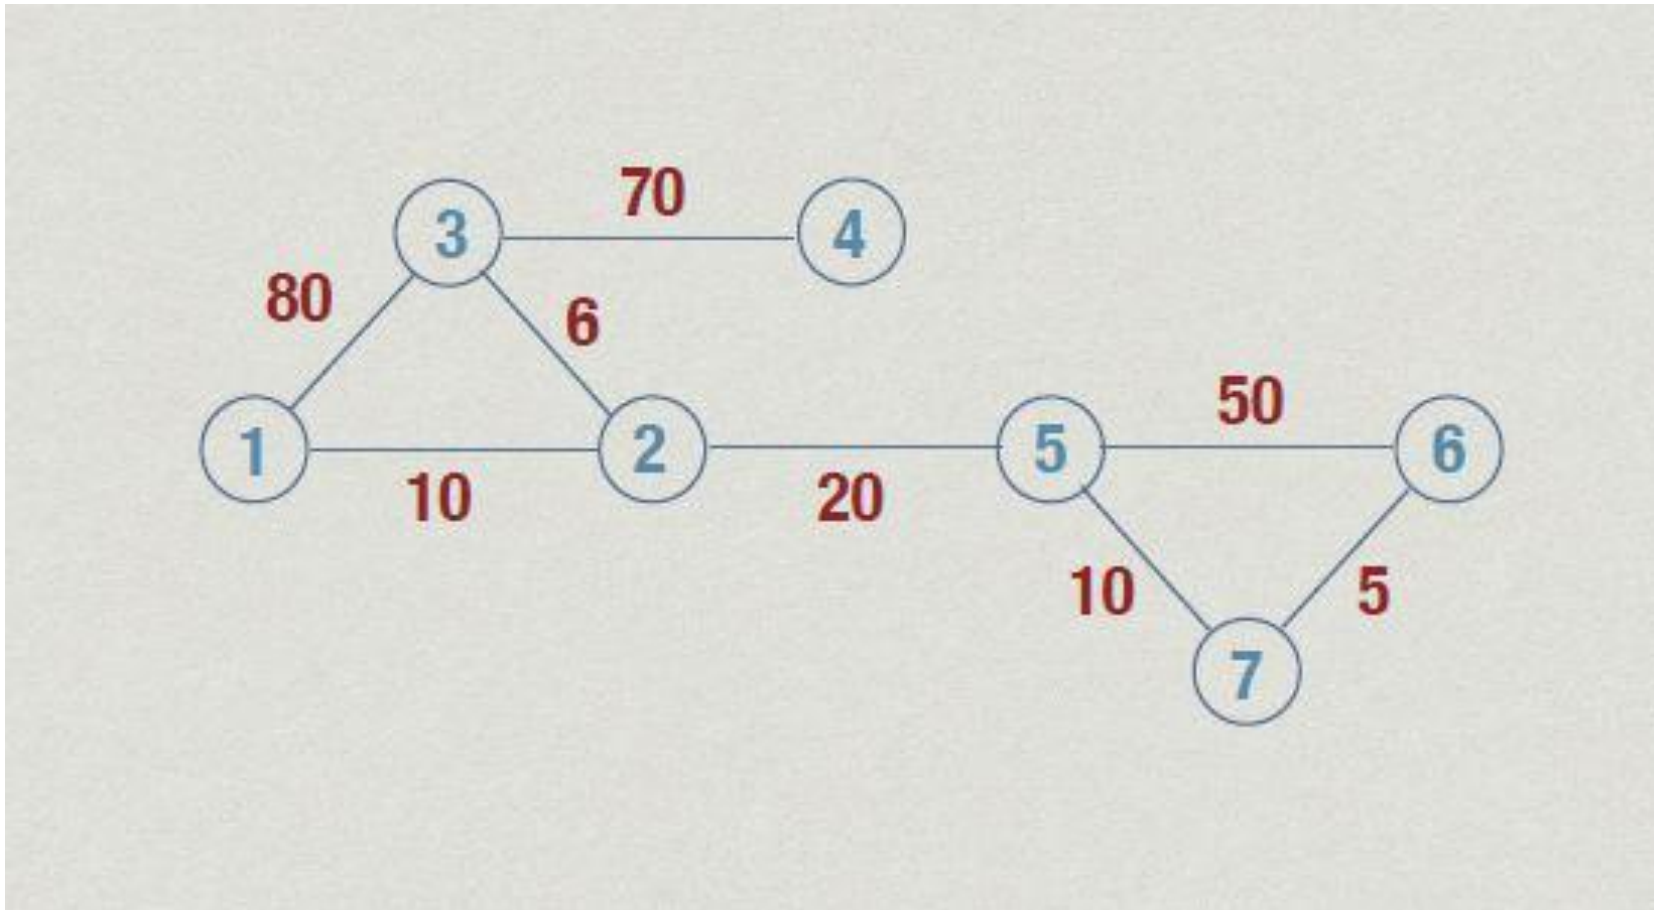
Single source shortest paths ...



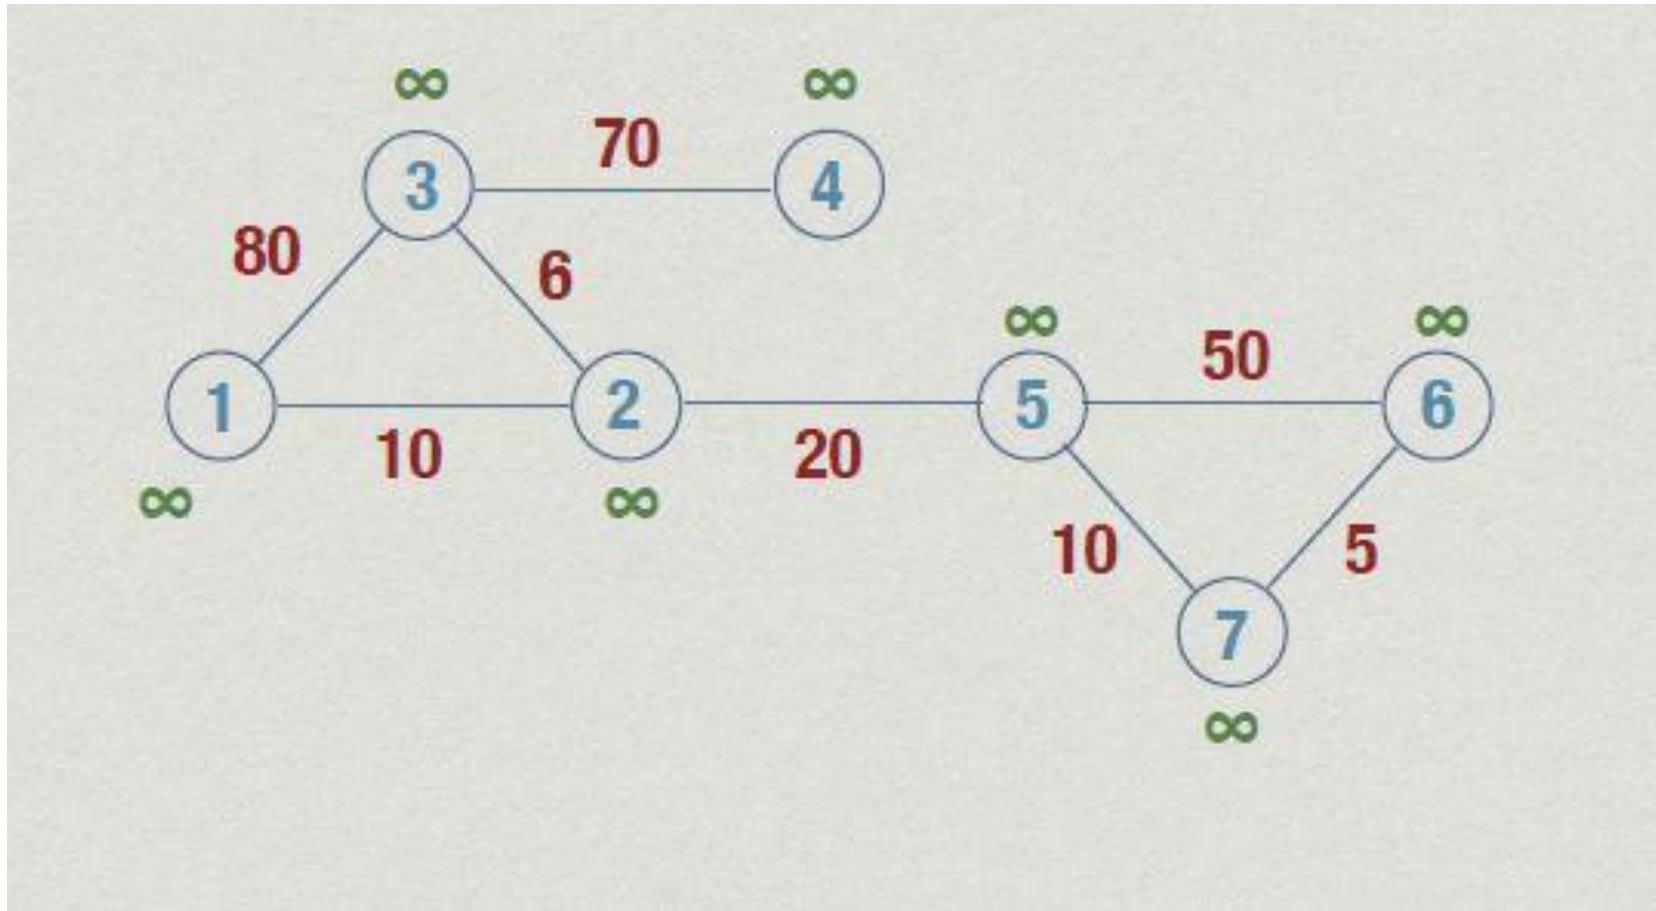
Single source shortest paths ...

- Compute expected time to burn of each vertex
- Update this each time a new vertex burns

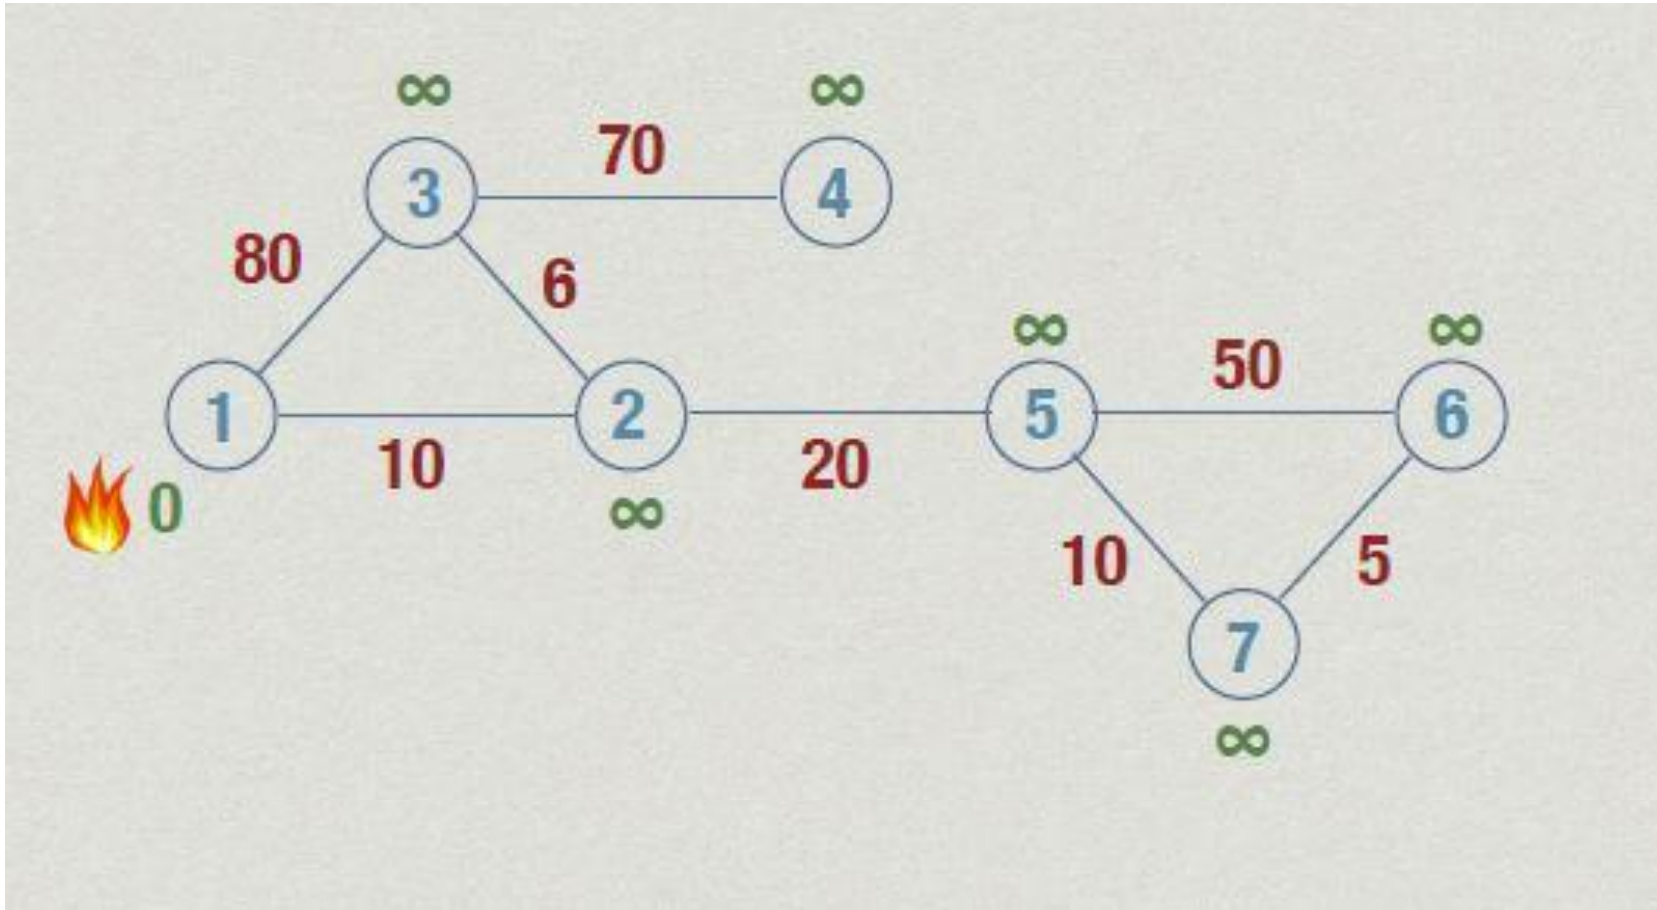
Single source shortest paths ...



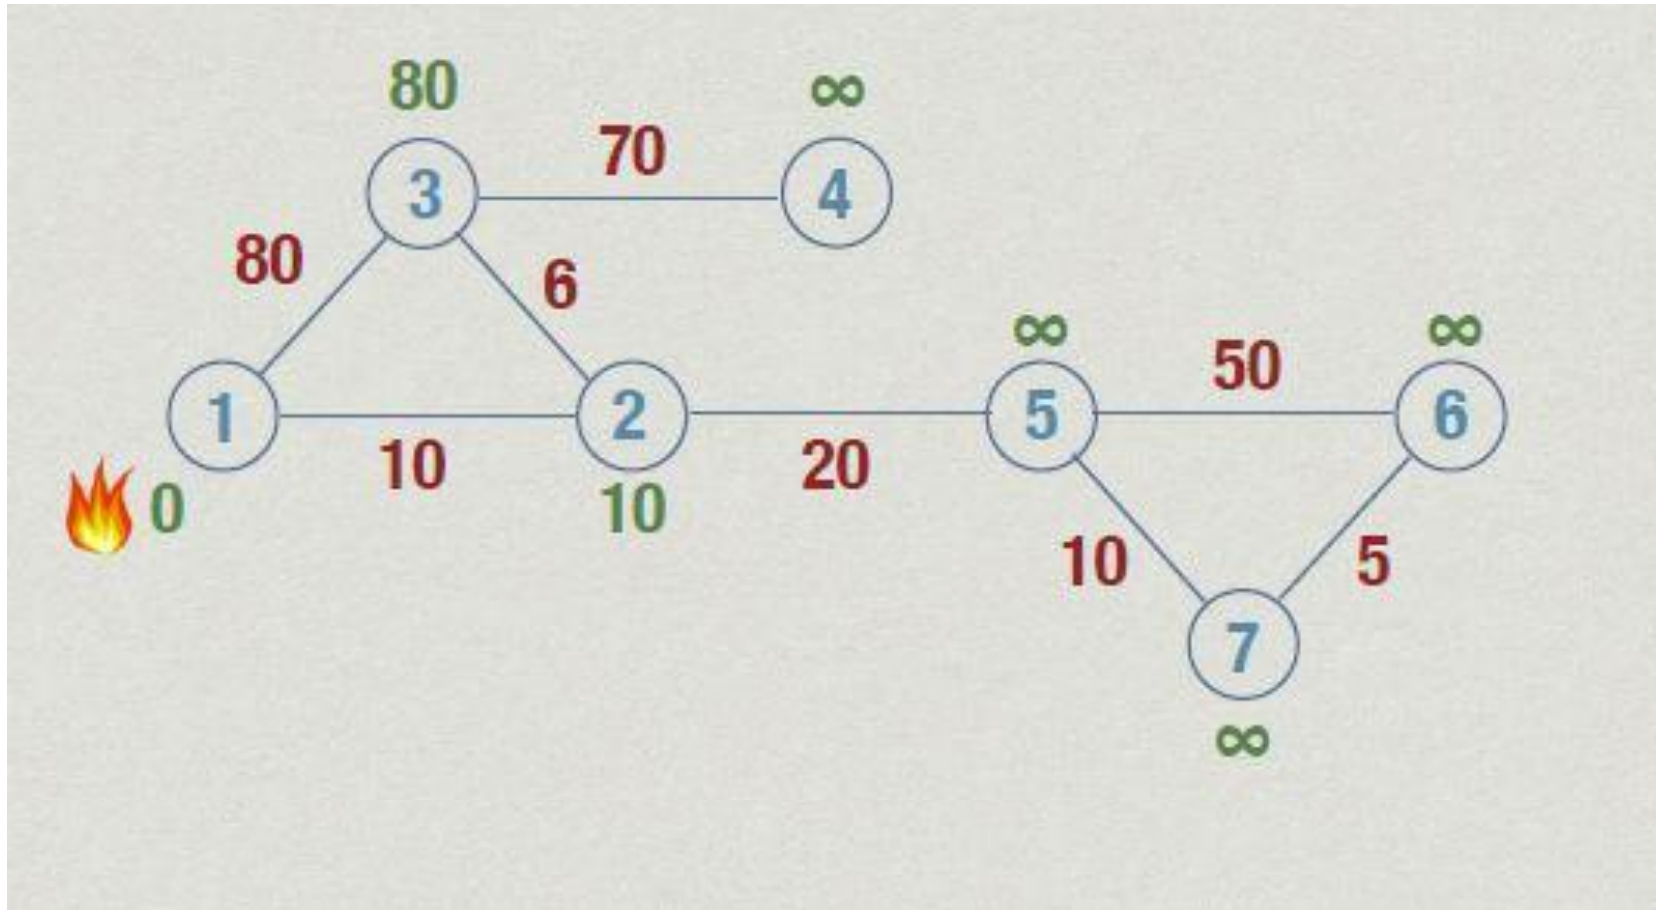
Single source shortest paths ...



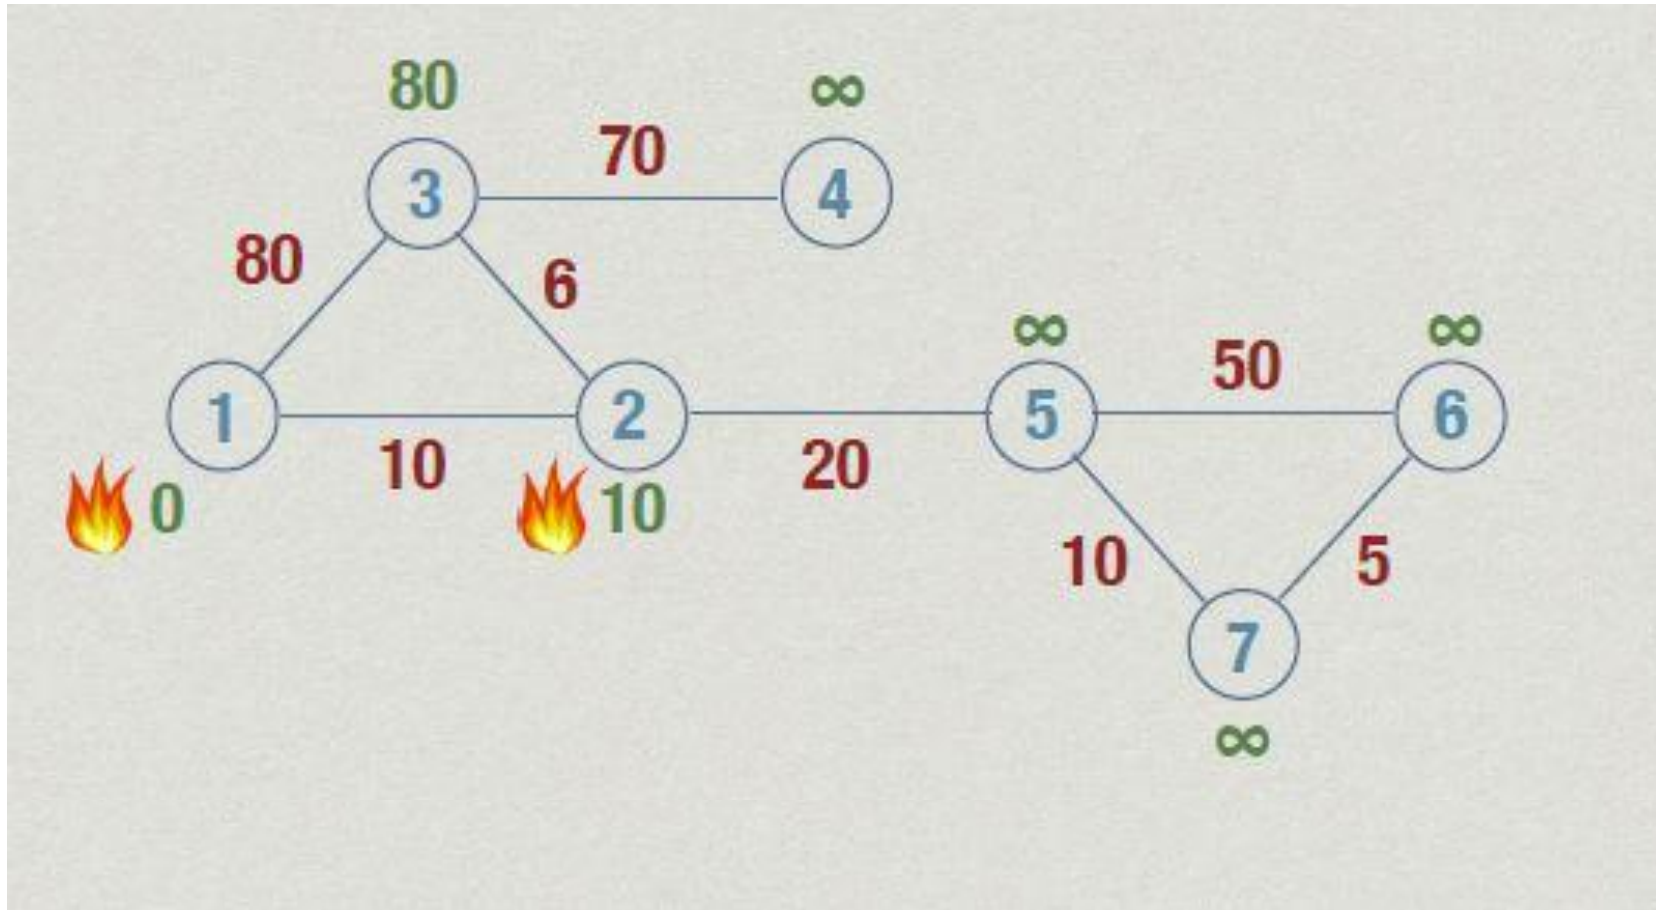
Single source shortest paths ...



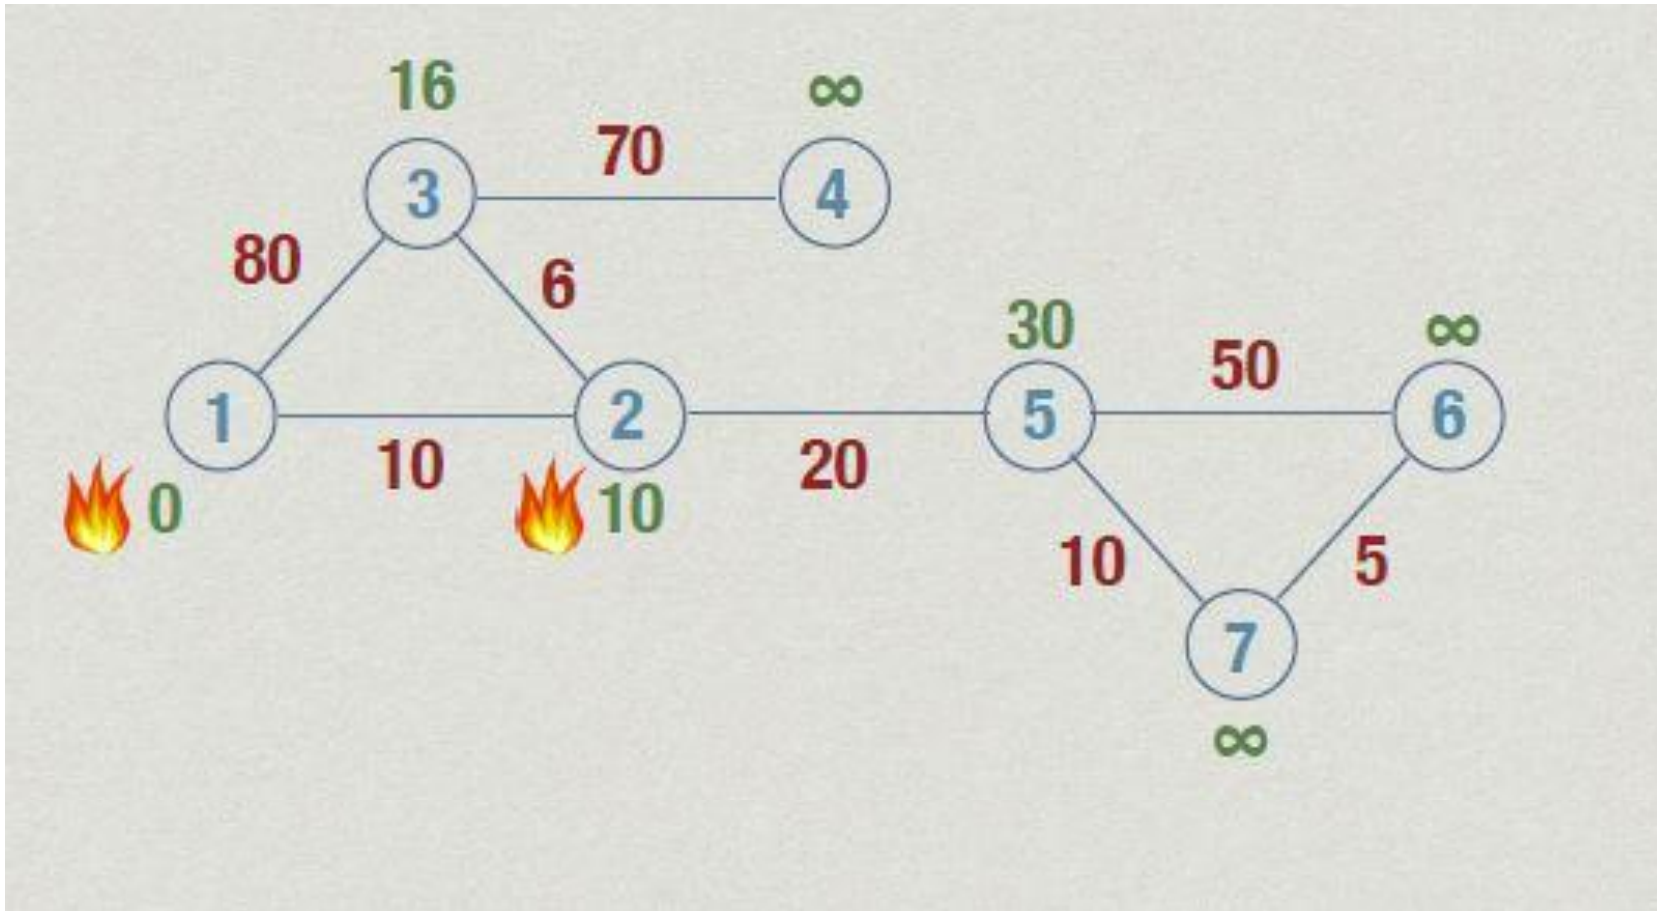
Single source shortest paths ...



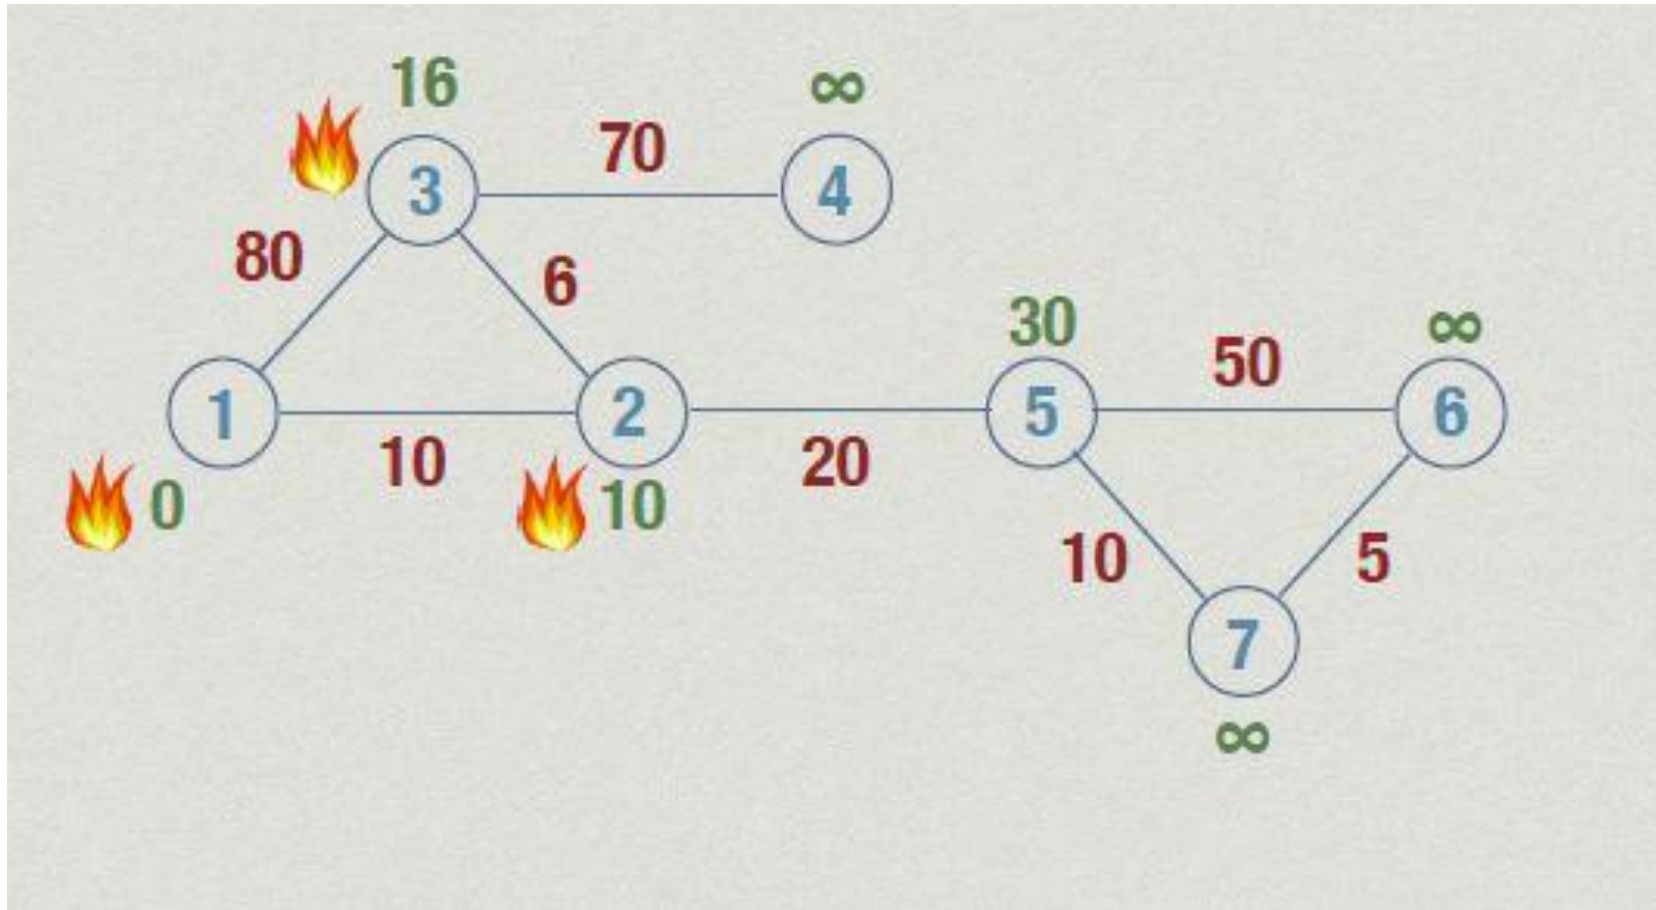
Single source shortest paths ...



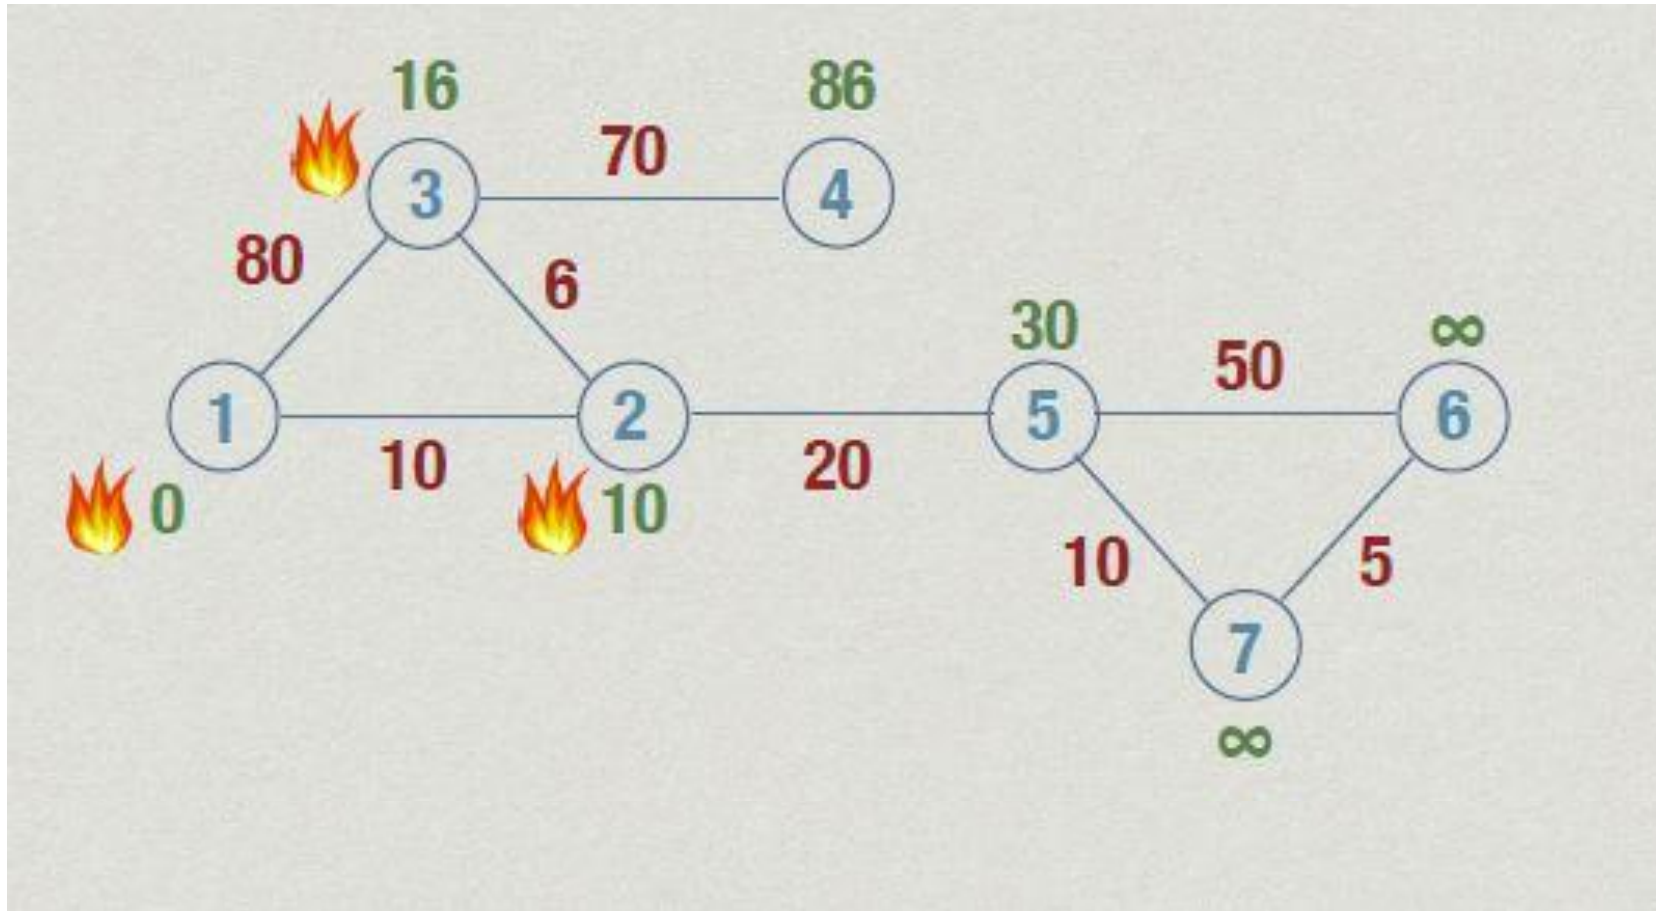
Single source shortest paths ...



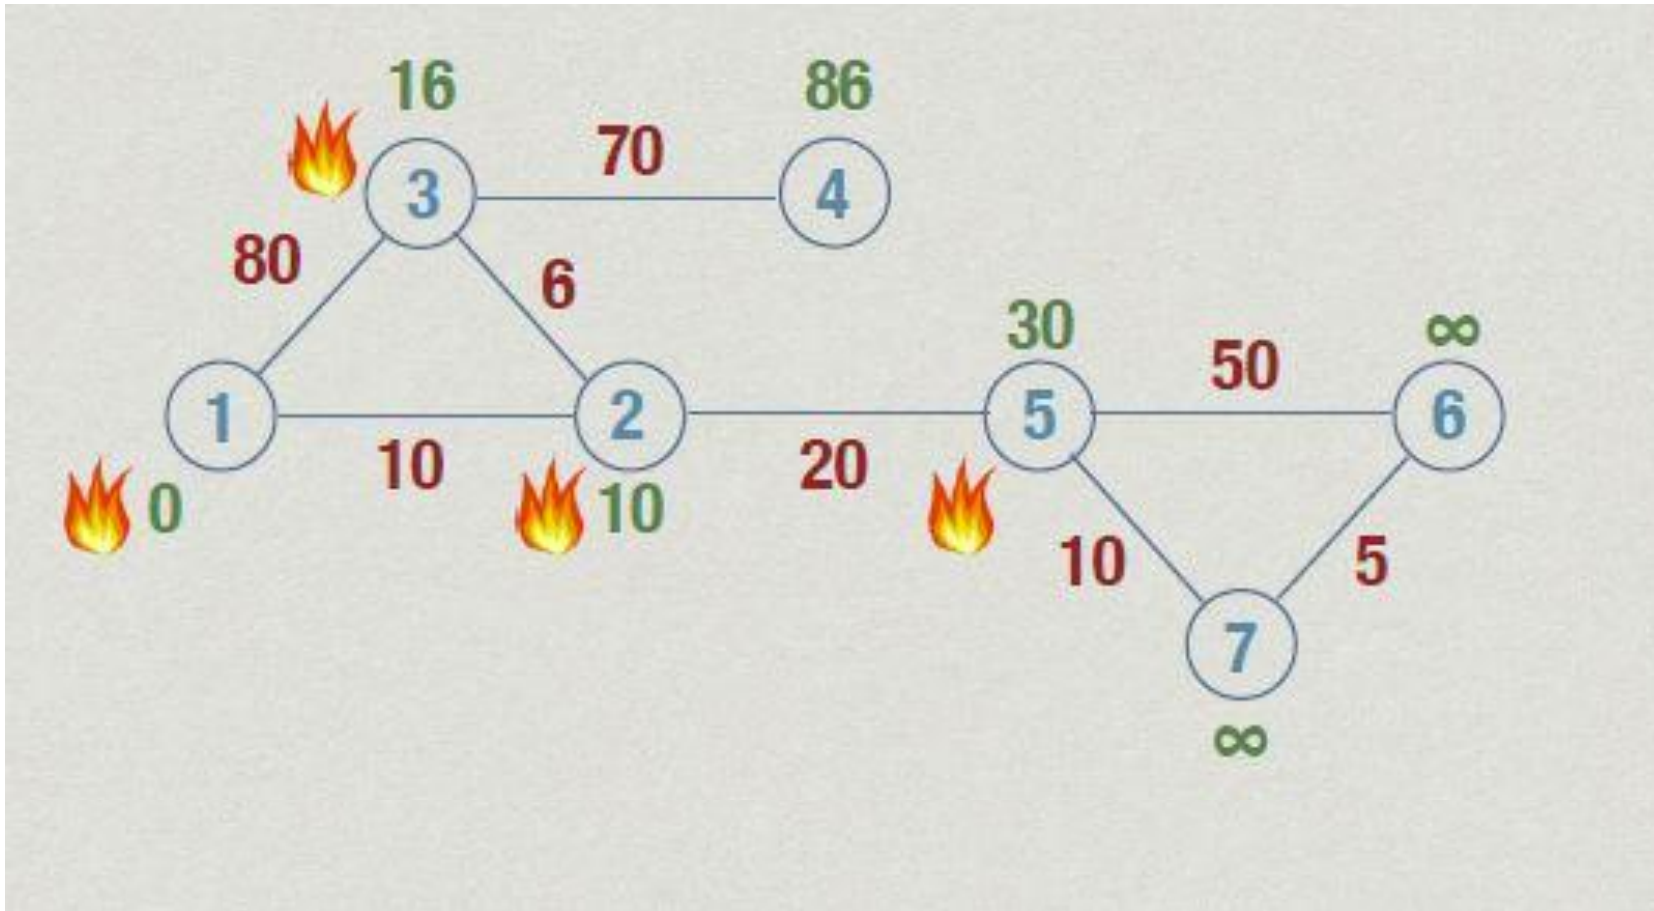
Single source shortest paths ...



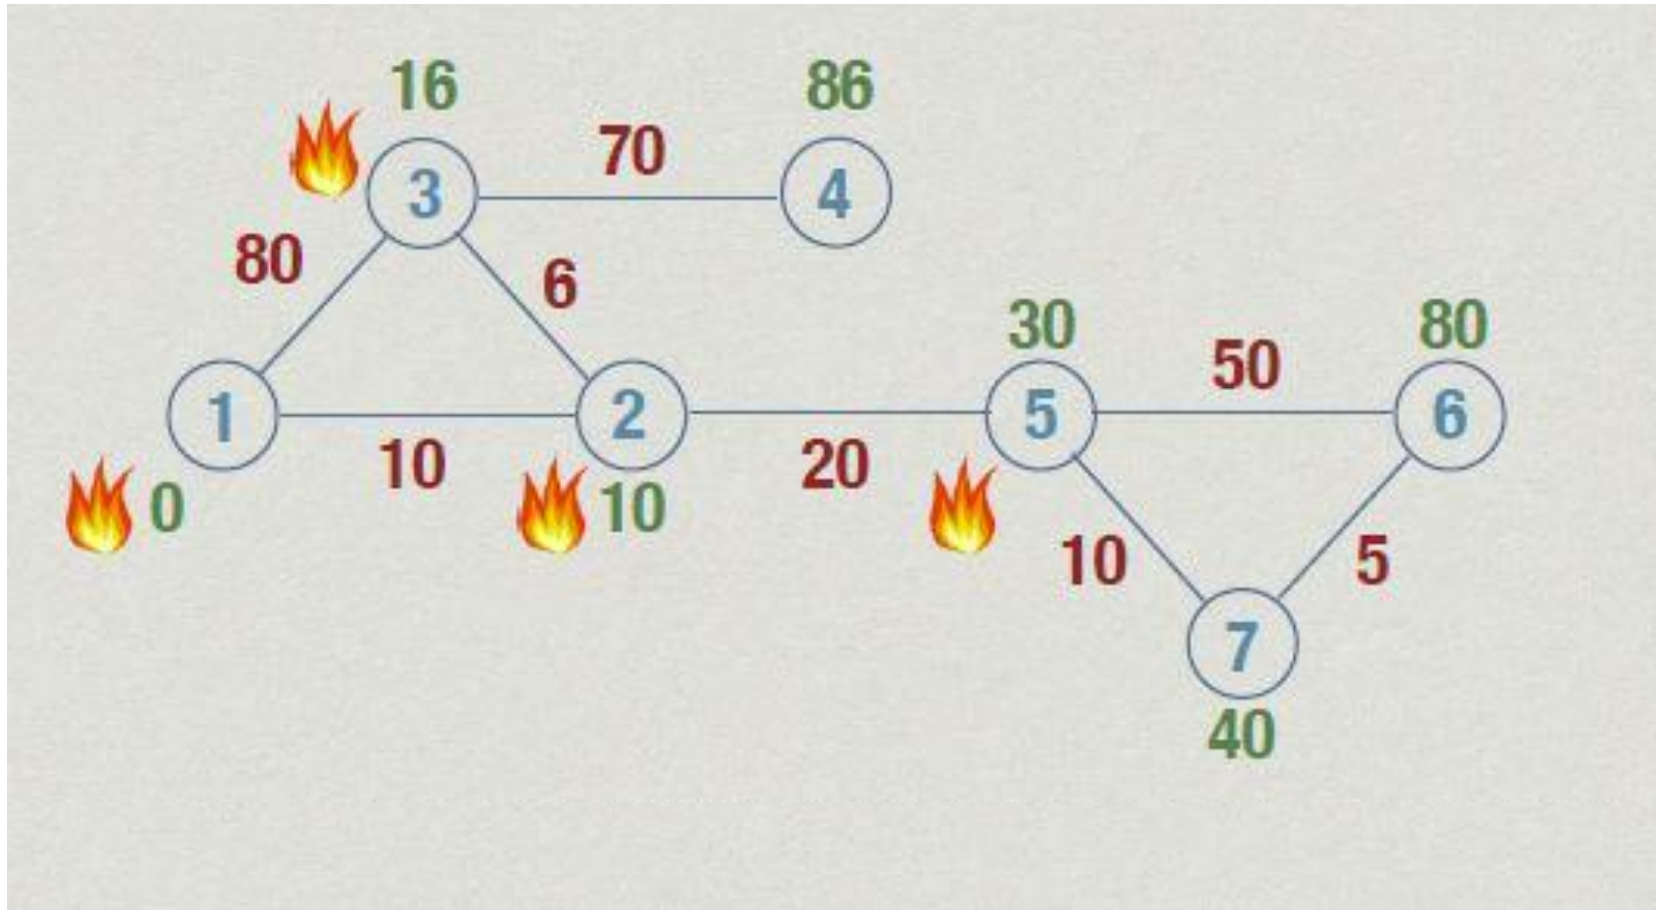
Single source shortest paths ...



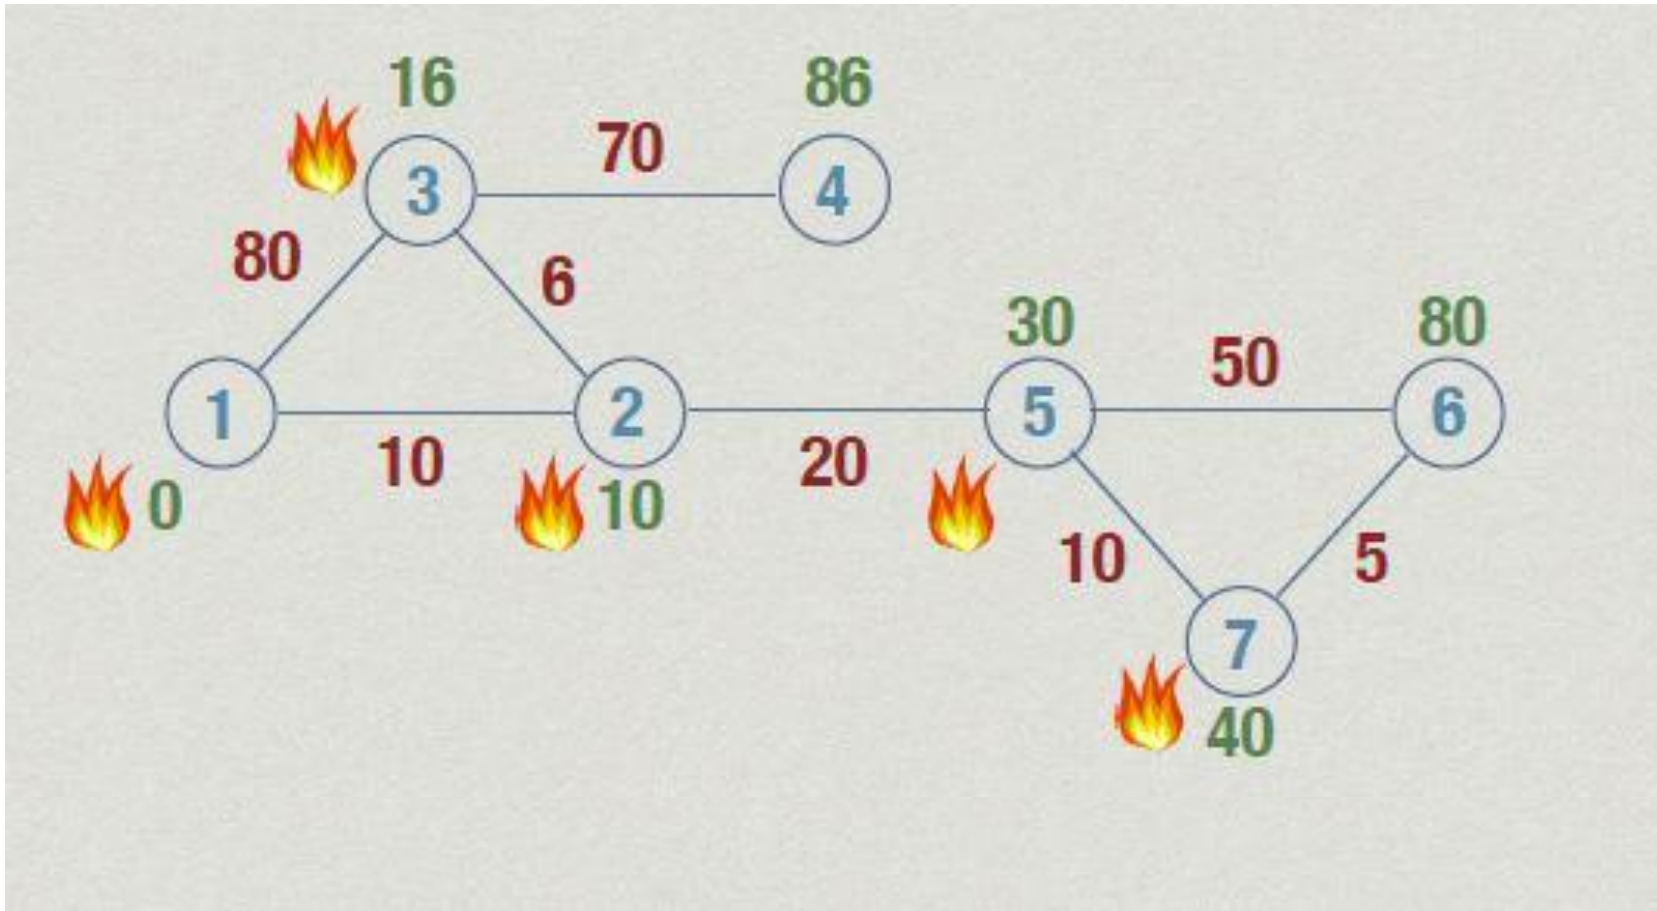
Single source shortest paths ...



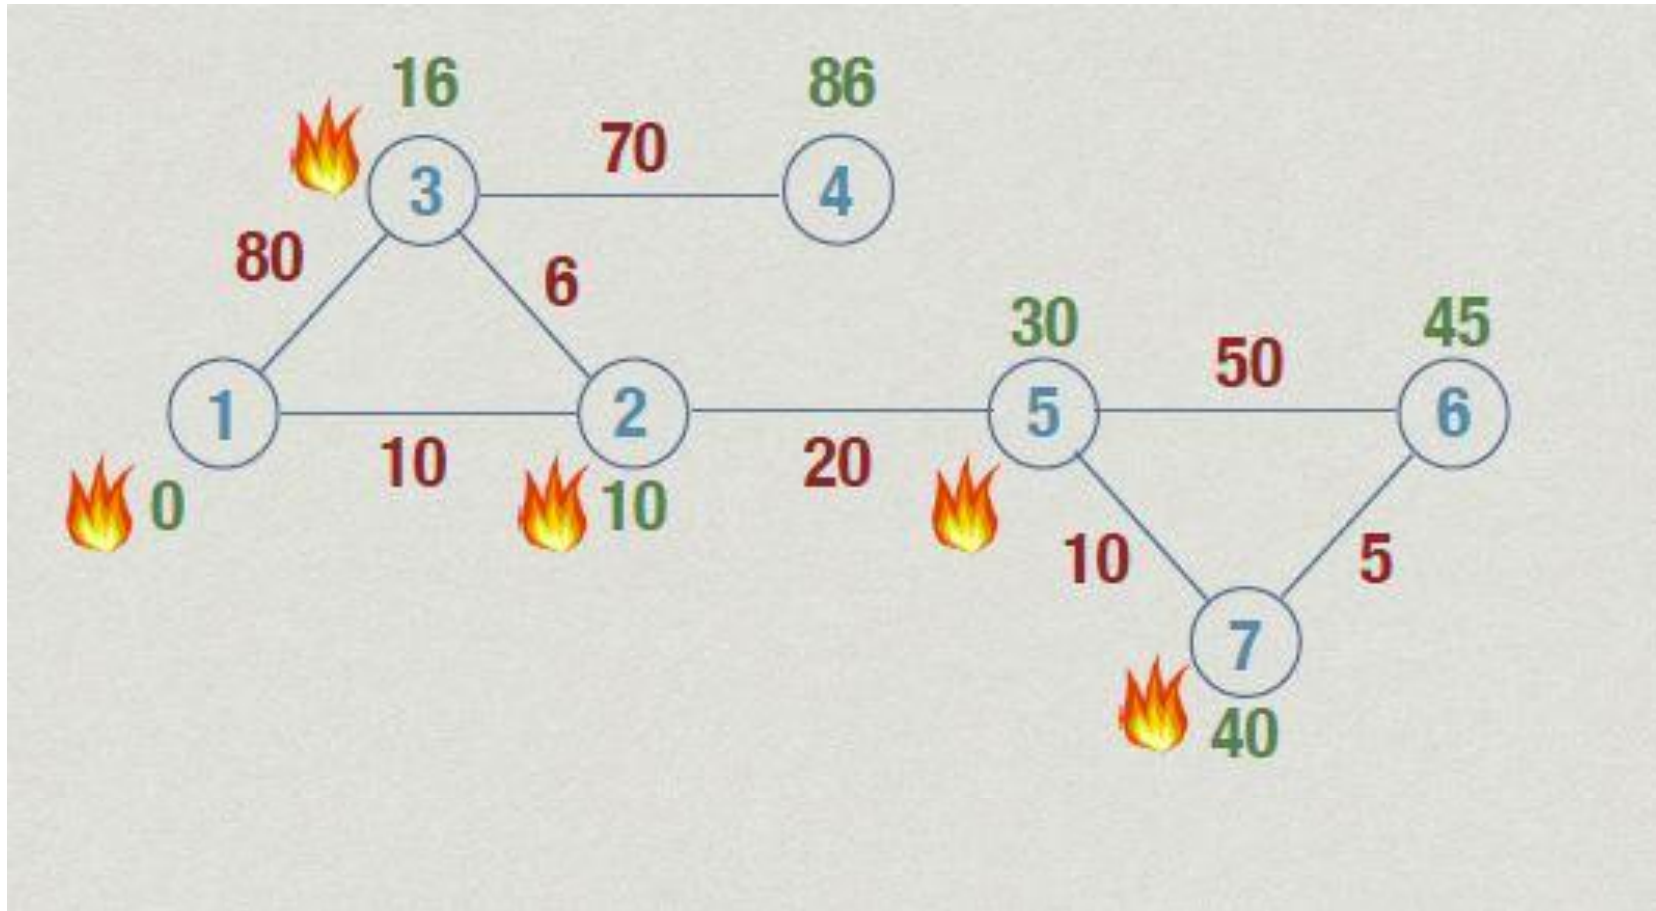
Single source shortest paths ...



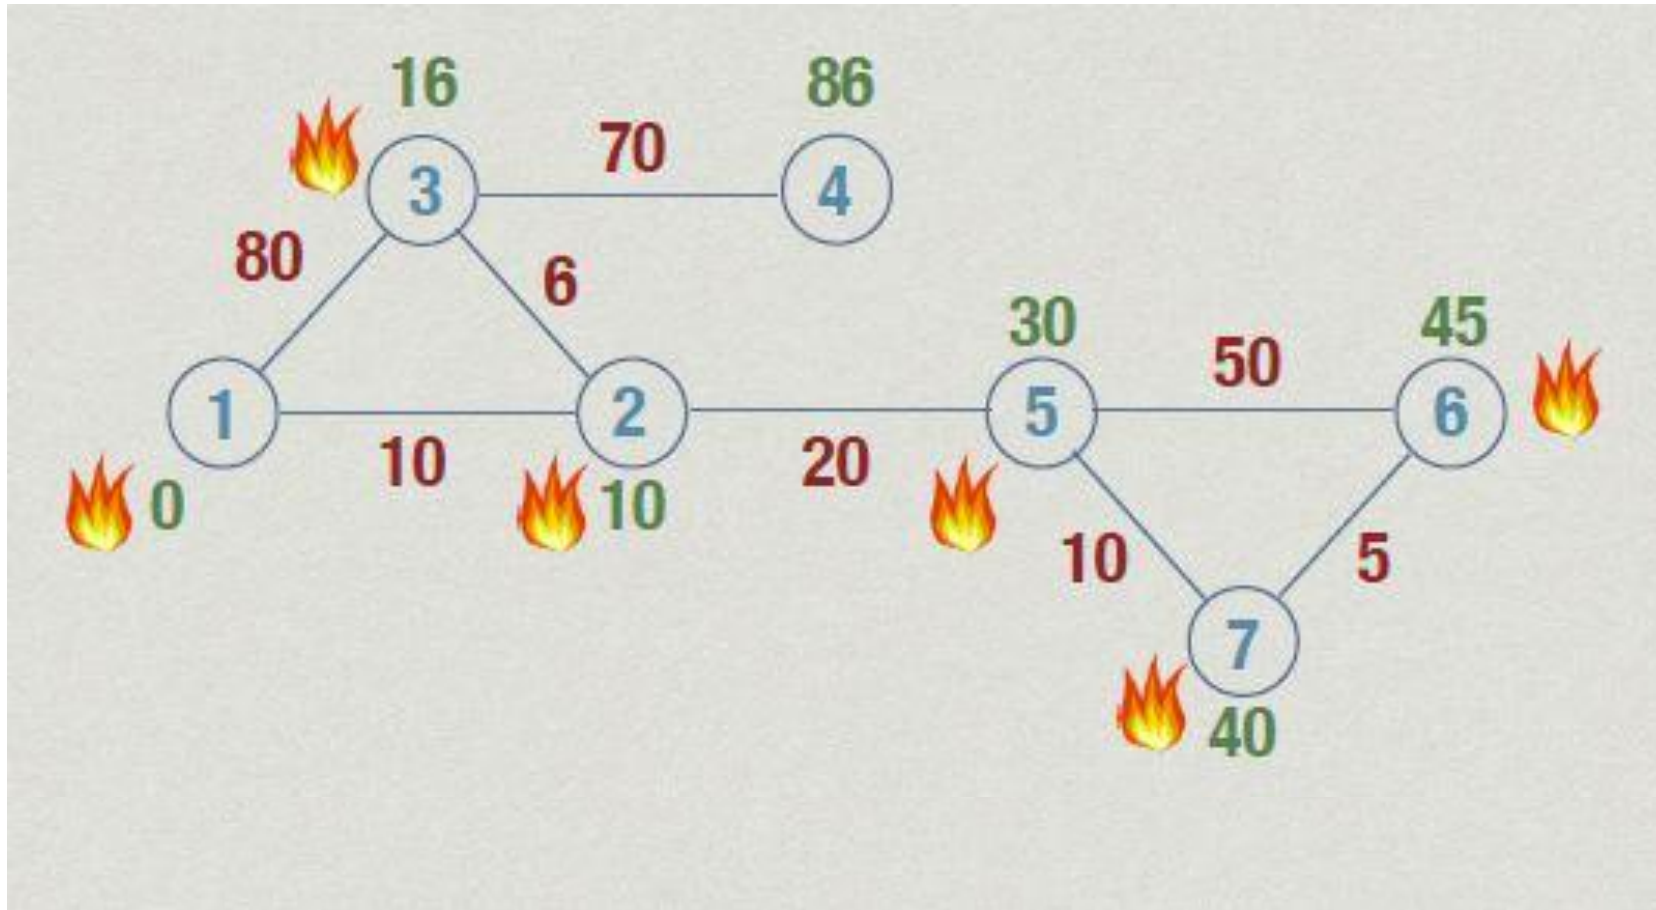
Single source shortest paths ...



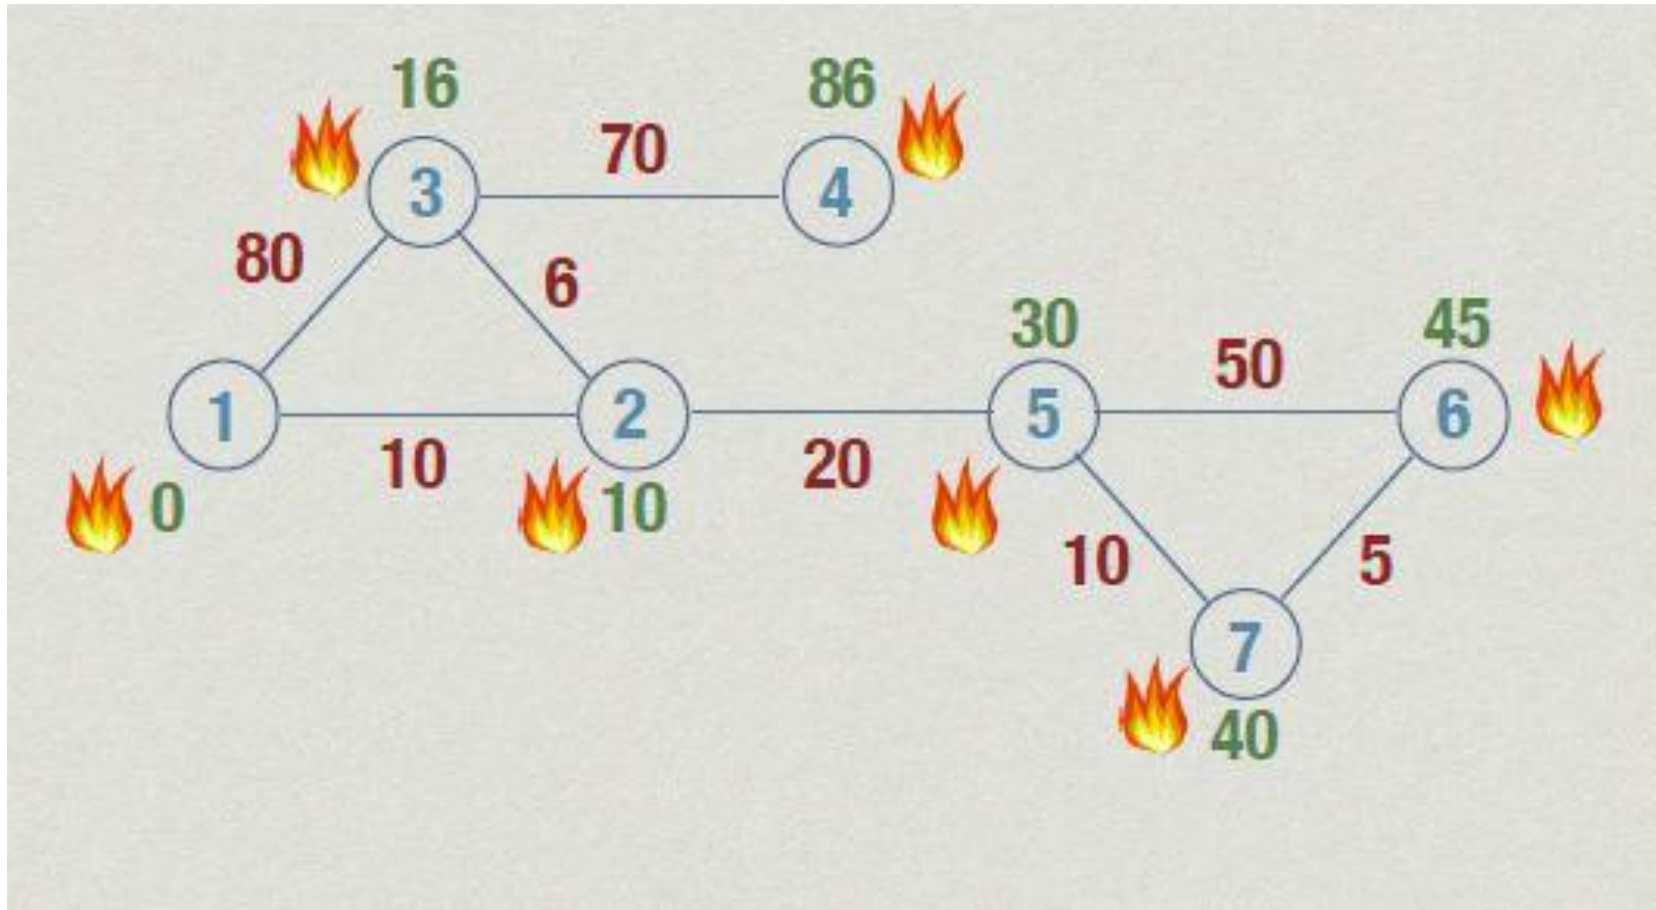
Single source shortest paths ...



Single source shortest paths ...



Single source shortest paths ...



Single source shortest paths ...

- Algorithmically
- Maintain two arrays
 - *BurnVertices*[], initially *False* for all *i*
 - *ExpectedBurnTime*[], initially ∞ for all *i*

Single source shortest paths ...

- Algorithmically
- Maintain two arrays
 - *BurnVertices*[], initially *False* for all i
 - *ExpectedBurnTime*[], initially ∞ for all i
 - For ∞ , use sum of all edge weights +1

Single source shortest paths ...

- Algorithmically
- Maintain two arrays
 - *BurnVertices*[], initially *False* for all *i*
 - *ExpectedBurnTime*[], initially ∞ for all *i*
 - For ∞ , use sum of all edge weights +1
 - Set *ExpectedBurnTime*[1] = 0
 - Repeat, until all vertices are burnt
 - Find *j* with minimum *ExpectedBurnTime*
 - Set *BurnVertices*[*j*] = *True*
 - Recompute *ExpectedBurnTime*[*k*] for each neighbor *k* of *j*

Dijkstra's algorithm

```
function ShortestPaths(s){ // assume source is s
    for i = 1 to n
        BV[i] = False; EBT[i] = infinity
    EBT[s] = 0
    for i = 1 to n
        Choose u such that BV[u] == False and EBT[u] is minimum
        BV[u] = True
        for each edge (u, v) with BV[v] == False
            if EBT[v] > EBT[u] + weight(u, v)
                EBT[v] = EBT[u] + weight(u, v)
```

Dijkstra's algorithm ...

```
function ShortestPaths(s){ // assume source is s
    for i = 1 to n
        Visited[i] = False; Distance[i] = infinity
    Distance[s] = 0
    for i = 1 to n
        Choose u such that Visited[u] == False and Distance[u] is
minimum
        Visited[u] = True
        for each edge (u, v) with Visited[v] == False
            if Distance[v] > Distance[u] + weight(u, v)
                Distance[v] = Distance[u] + weight(u, v)
```

Dijkstra's algorithm: Analysis

Dijkstra's algorithm: Analysis

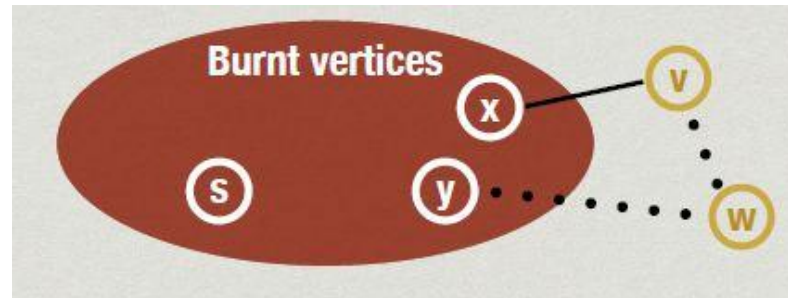
- Dijkstra's algorithm is greedy
 - Select vertex with minimum expected burn time

Dijkstra's algorithm: Analysis

- Dijkstra's algorithm is greedy
 - Select vertex with minimum expected burn time
- Need to prove that greedy strategy is optimal
- Most times, greedy approach fails
 - Current best may not be globally optimal

Dijkstra's algorithm: Analysis ...

- Correctness
- Each new shortest path we discover extends an earlier one
- By induction, assume we have identified shortest paths to all vertices already burnt



- Next vertex to burn is v , via x
- Can't later find a shorter path from y to w to v

Dijkstra's algorithm: Analysis ...

- Complexity: ?

Dijkstra's algorithm: Analysis ...

- Complexity
- Outer loop runs n times
 - In each iteration, we burn one vertex
 - $O(n)$ scan to find minimum burn time vertex
- Each time we burn a vertex v , we have to scan all its neighbors to update burn times
 - $O(n)$ scan of adjacency matrix to find all neighbors
- Overall $O(n^2)$

Dijkstra's algorithm: Analysis ...

- Complexity
- Does adjacency list help?

Dijkstra's algorithm: Analysis ...

- Complexity
- Does adjacency list help?
 - Scan neighbors to update burn times
 - $O(m)$ across all iterations
- However, identifying minimum burn time vertex still takes $O(n)$ in each iteration
- Still $O(n^2)$

Dijkstra's algorithm: Analysis ...

- Complexity
- Can maintain ExpectedBurnTime in a more sophisticated data structure?

Dijkstra's algorithm: Analysis ...

- Complexity
- Can maintain ExpectedBurnTime in a more sophisticated data structure?
 - Different types of trees (heaps, red-black trees) allow both the following in $O(\log n)$ time
 - Find and delete minimum
 - Insert or update value

Dijkstra's algorithm: Analysis ...

- Complexity
- Can maintain ExpectedBurnTime in a more sophisticated data structure?
 - Different types of trees (heaps, red-black trees) allow both the following in $O(\log n)$ time
 - Find and delete minimum
 - Insert or update value
- With such a tree
 - Finding minimum burn time vertex takes $O(\log n)$
 - With adjacency list, updating burn times take $O(\log n)$ each, total $O(m)$ edges
- Overall $O(n \log n + m \log n) = O((n + m) \log n)$

Dijkstra's algorithm: Limitations

- What if edge weights can be negative?

Dijkstra's algorithm: Limitations

- What if edge weights can be negative?
- Our correctness argument is no longer valid

Dijkstra's algorithm: Limitations ...

- Why negative weights?

Dijkstra's algorithm: Limitations ...

- Why negative weights?
- Weight represent money
 - Taxi driver earns money from airport to city, travels empty to next pick-up point
 - Some segment earn money. Some lose money
- Chemistry
 - Nodes are compounds, edges are reactions
 - Weights are energy absorbed/released by reaction

Negative weights ...

- **Negative cycle:** loop with a negative weight
 - Problem is not well defined with negative cycles
 - Repeatedly traversing cycle pushes down cost
- With negative edges, but no negative cycles, shortest paths do exist

About shortest paths

- Shortest paths will never loop
 - Never visit the same vertex twice
 - At most length $n - 1$

About shortest paths

- Shortest paths will never loop
 - Never visit the same vertex twice
 - At most length $n - 1$
- Every prefix of a shortest path is itself a shortest path
 - Suppose the shortest path from s to t is
$$s \rightarrow v_1 \rightarrow v_2 \rightarrow v_3 \dots \rightarrow v_m \rightarrow t$$
 - Every prefix $s \rightarrow v_1 \rightarrow v_2 \rightarrow v_3 \dots \rightarrow v_r$ is a shortest path to v_r

Updating Distance()

- When vertex j is “burnt”, for each (j, k) update
 - $Distance(k) = \min(Distance(k), Distance(j) + weight(j, k))$
- Refer to this as $update(j, k)$
- Dijkstra’s algorithm
 - When we compute $update(j, k)$, $Distance(j)$ is always guaranteed to be correct distance to j
- What we can say in general?

Properties of $\text{update}(j, k)$

- $\text{update}(j, k)$:
 - $\text{Distance}(k) = \min(\text{Distance}(k), \text{Distance}(j) + \text{weight}(j, k))$
- $\text{Distance}(k)$ is no more than $\text{Distance}(j) + \text{weight}(j, k)$
- If $\text{Distance}(j)$ is correct and j is the second-last node on shortest path to k , $\text{Distance}(k)$ is correct
- Update is safe
 - $\text{Distance}(k)$ never becomes “too small”
 - Redundant updates cannot hurt

Updating distance() ...

- $update(j, k)$
 - $Distance(k) = \min(Distance(k), Distance(j) + weight(j, k))$
- Dijkstra's algorithm performs a particular “greedy” sequence of updates
 - Computes shortest paths without negative weights
- With negative edges, this sequence does not work
- Is there some sequence that dose work?

Updating distance() ...

- Suppose the shortest path from s to t is

$$s \rightarrow v_1 \rightarrow v_2 \rightarrow v_3 \dots \rightarrow v_m \rightarrow t$$

- If our update sequence includes ..., $update(s, v_1)$, ..., $update(v_1, v_2)$, ..., $update(v_2, v_3)$, ..., $update(v_m, t)$, ..., in that order, $Distance(t)$ will be computed correctly
- If $Distance(j)$ is correct and j is second-last node on shortest path to k , $Distance(k)$ is correct after $update(j, k)$

Bellman-Ford algorithm

- Initialize $distance(s) = 0, distance(u) = \infty$ for all other vertices
- Update all edges $n - 1$ times!

Bellman-Ford algorithm

- Initialize $distance(s) = 0, distance(u) = \infty$ for all other vertices
- Update all edges $n - 1$ times!

Iteration 1
...
update(s, v ₁)
...
update(v ₁ , v ₂)
...
update(v ₂ , v ₃)
...
update(v _m , t)
...

Bellman-Ford algorithm

- Initialize $distance(s) = 0, distance(u) = \infty$ for all other vertices
- Update all edges $n - 1$ times!

Iteration 1	Iteration 2
...	...
update(s,v ₁)	update(s,v ₁)
...	...
update(v ₁ ,v ₂)	update(v ₁ ,v ₂)
...	...
update(v ₂ ,v ₃)	update(v ₂ ,v ₃)
...	...
update(v _m ,t)	update(v _m ,t)
...	...

Bellman-Ford algorithm

- Initialize $distance(s) = 0, distance(u) = \infty$ for all other vertices
- Update all edges $n - 1$ times!

Iteration 1	Iteration 2	...
...
update(s,v ₁)	update(s,v ₁)	...
...
update(v ₁ ,v ₂)	update(v ₁ ,v ₂)	...
...
update(v ₂ ,v ₃)	update(v ₂ ,v ₃)	...
...
update(v _m ,t)	update(v _m ,t)	...
...

Bellman-Ford algorithm

- Initialize $distance(s) = 0, distance(u) = \infty$ for all other vertices
- Update all edges $n - 1$ times!

Iteration 1	Iteration 2	...	Iteration n-1
...
update(s,v ₁)	update(s,v ₁)	...	update(s,v ₁)
...
update(v ₁ ,v ₂)	update(v ₁ ,v ₂)	...	update(v ₁ ,v ₂)
...
update(v ₂ ,v ₃)	update(v ₂ ,v ₃)	...	update(v ₂ ,v ₃)
...
update(v _m ,t)	update(v _m ,t)	...	update(v _m ,t)
...

Bellman-Ford algorithm

- Initialize $distance(s) = 0, distance(u) = \infty$ for all other vertices
- Update all edges $n - 1$ times!

Iteration 1	Iteration 2	...	Iteration n-1
...
update(s,v ₁)	update(s,v ₁)	...	update(s,v ₁)
...
update(v ₁ ,v ₂)	update(v ₁ ,v ₂)	...	update(v ₁ ,v ₂)
...
update(v ₂ ,v ₃)	update(v ₂ ,v ₃)	...	update(v ₂ ,v ₃)
...
update(v _m ,t)	update(v _m ,t)	...	update(v _m ,t)
...

Bellman-Ford algorithm

- Initialize $distance(s) = 0, distance(u) = \infty$ for all other vertices
- Update all edges $n - 1$ times!

Iteration 1	Iteration 2	...	Iteration n-1
...
update(s,v ₁)	update(s,v ₁)	...	update(s,v ₁)
...
update(v ₁ ,v ₂)	update(v ₁ ,v ₂)	...	update(v ₁ ,v ₂)
...
update(v ₂ ,v ₃)	update(v ₂ ,v ₃)	...	update(v ₂ ,v ₃)
...
update(v _m ,t)	update(v _m ,t)	...	update(v _m ,t)
...

Bellman-Ford algorithm

- Initialize $distance(s) = 0, distance(u) = \infty$ for all other vertices
- Update all edges $n - 1$ times!

Iteration 1	Iteration 2	...	Iteration n-1
...
update(s,v ₁)	update(s,v ₁)	...	update(s,v ₁)
...
update(v ₁ ,v ₂)	update(v ₁ ,v ₂)	...	update(v ₁ ,v ₂)
...
update(v ₂ ,v ₃)	update(v ₂ ,v ₃)	...	update(v ₂ ,v ₃)
...
update(v _m ,t)	update(v _m ,t)	...	update(v _m ,t)
...

Bellman-Ford algorithm

function *BellmanFord*(*s*)//source *s*, with -ve weights

for *i* = 1 to *n*

Distance[*i*] = *infinity*

Distance[*s*] = 0

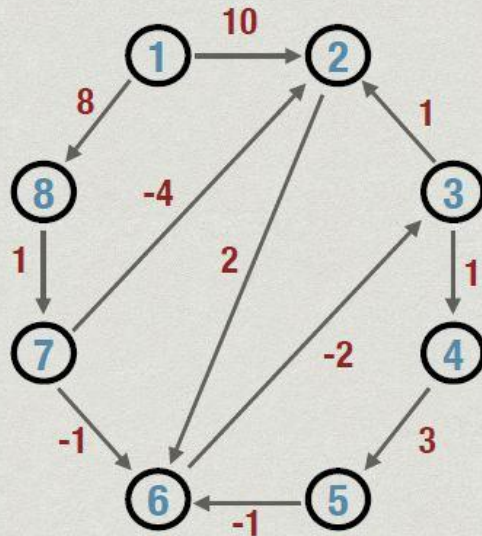
for *i* = 1 to *n* - 1 //repeat *n*-1 times

 for each *edge*(*j*, *k*) in *E*

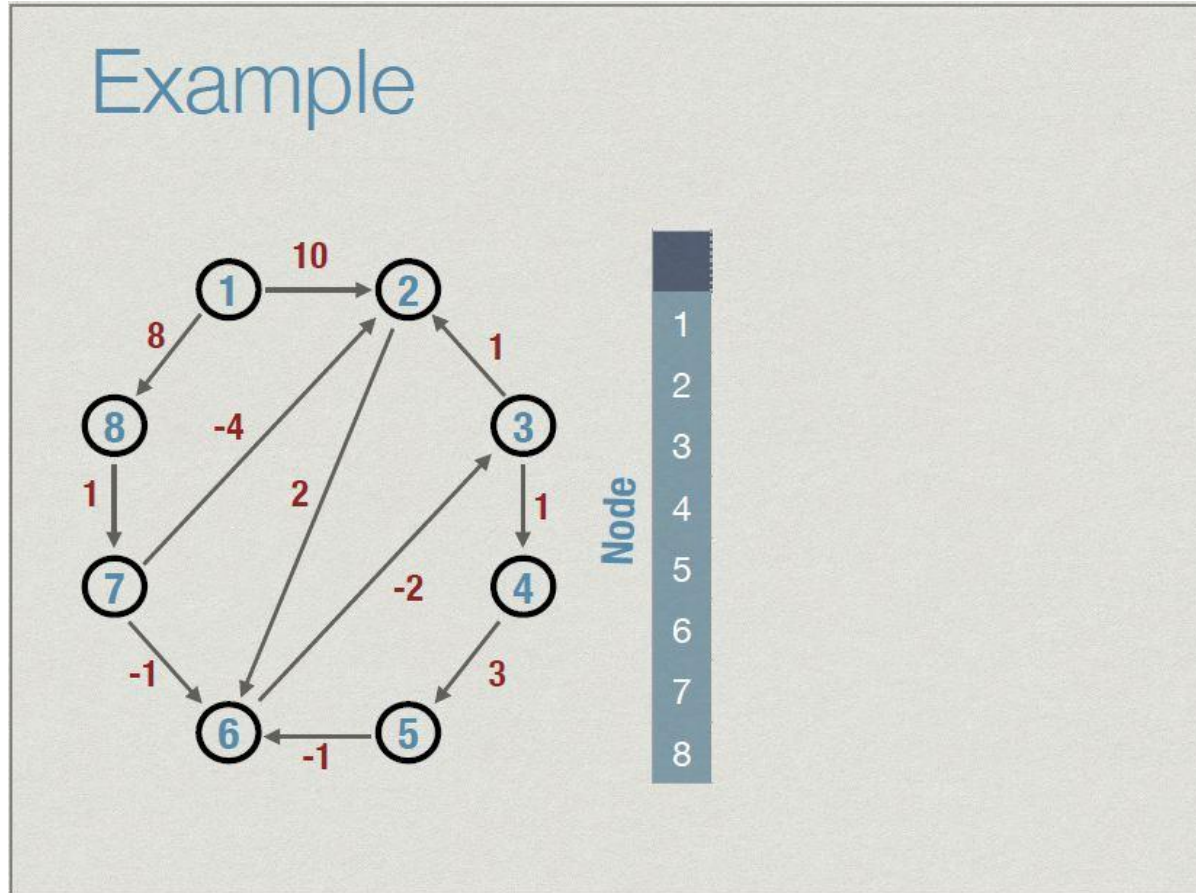
Distance(*k*) = min(*Distance*(*k*), *Distance*(*j*) +
 weight(*j*, *k*))

Bellman-Ford algorithm

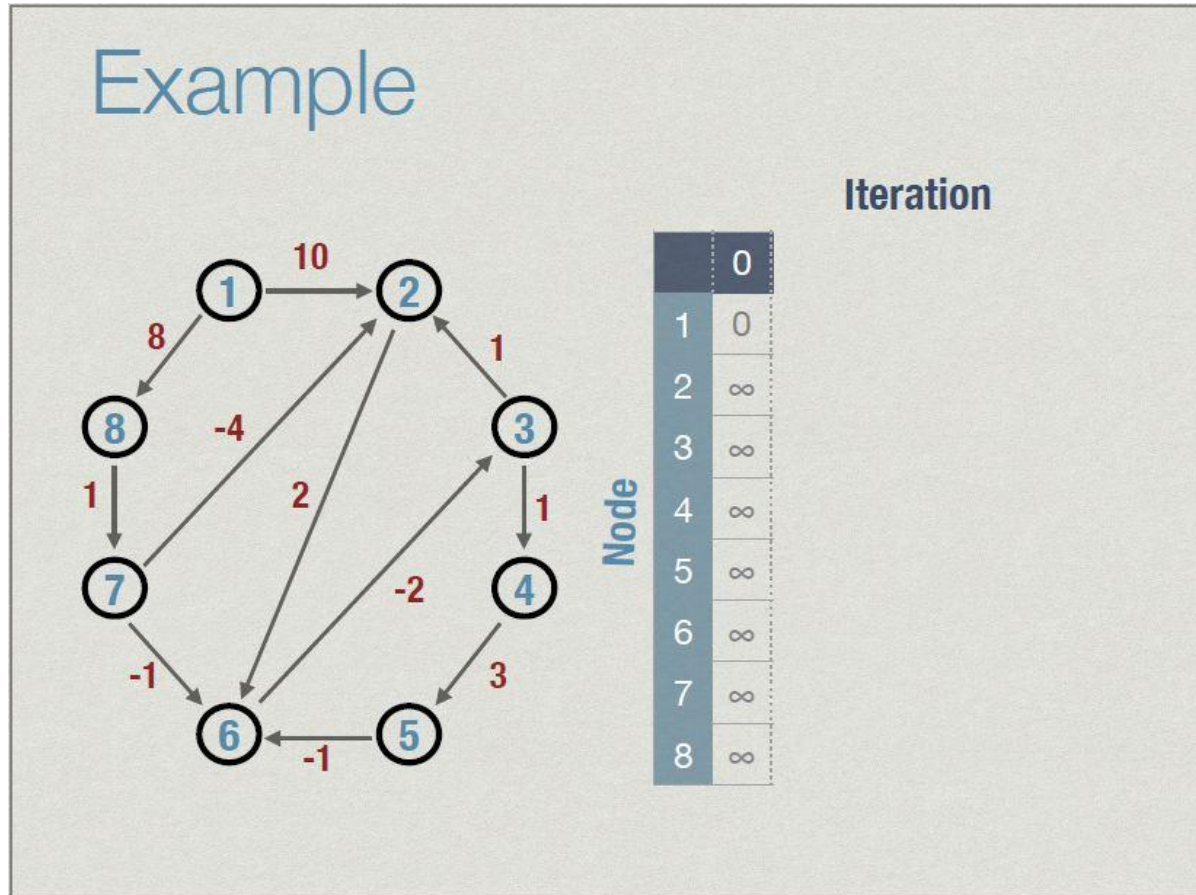
Example



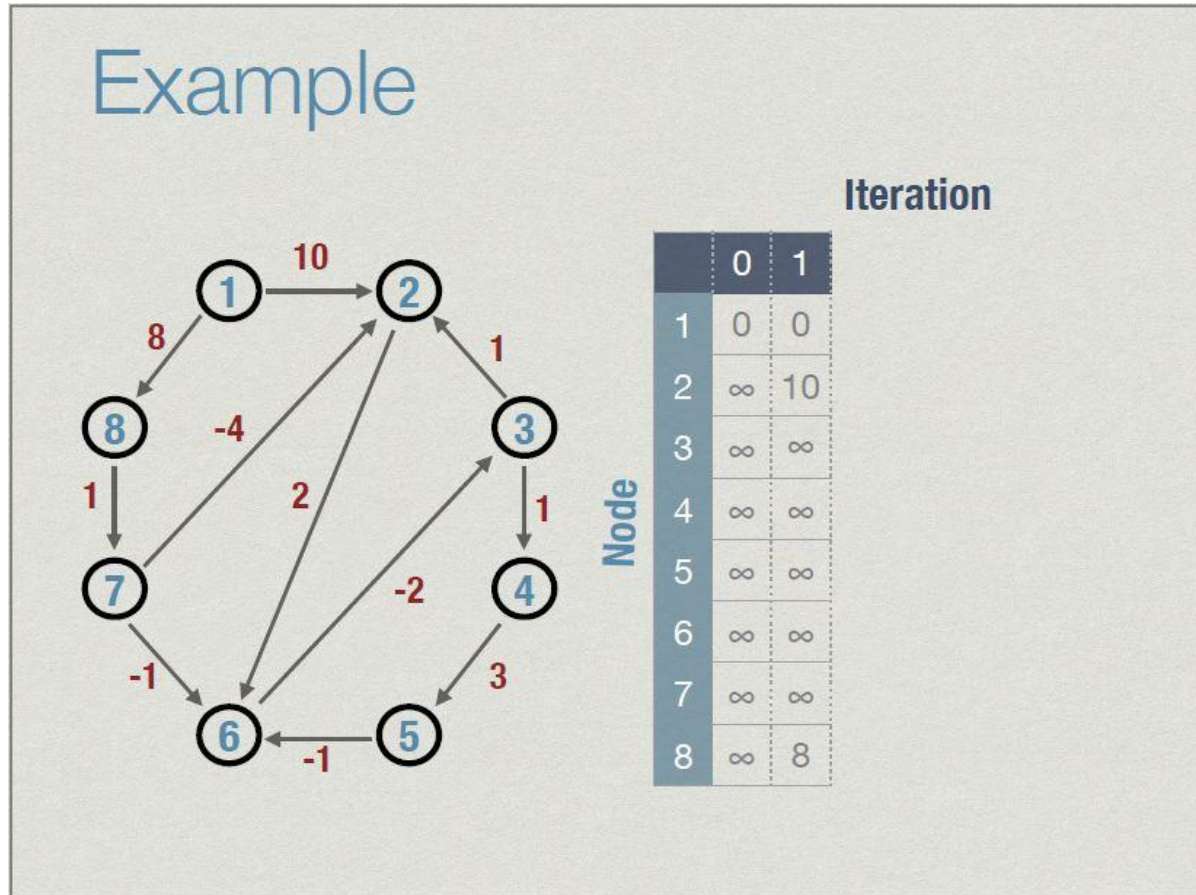
Bellman-Ford algorithm



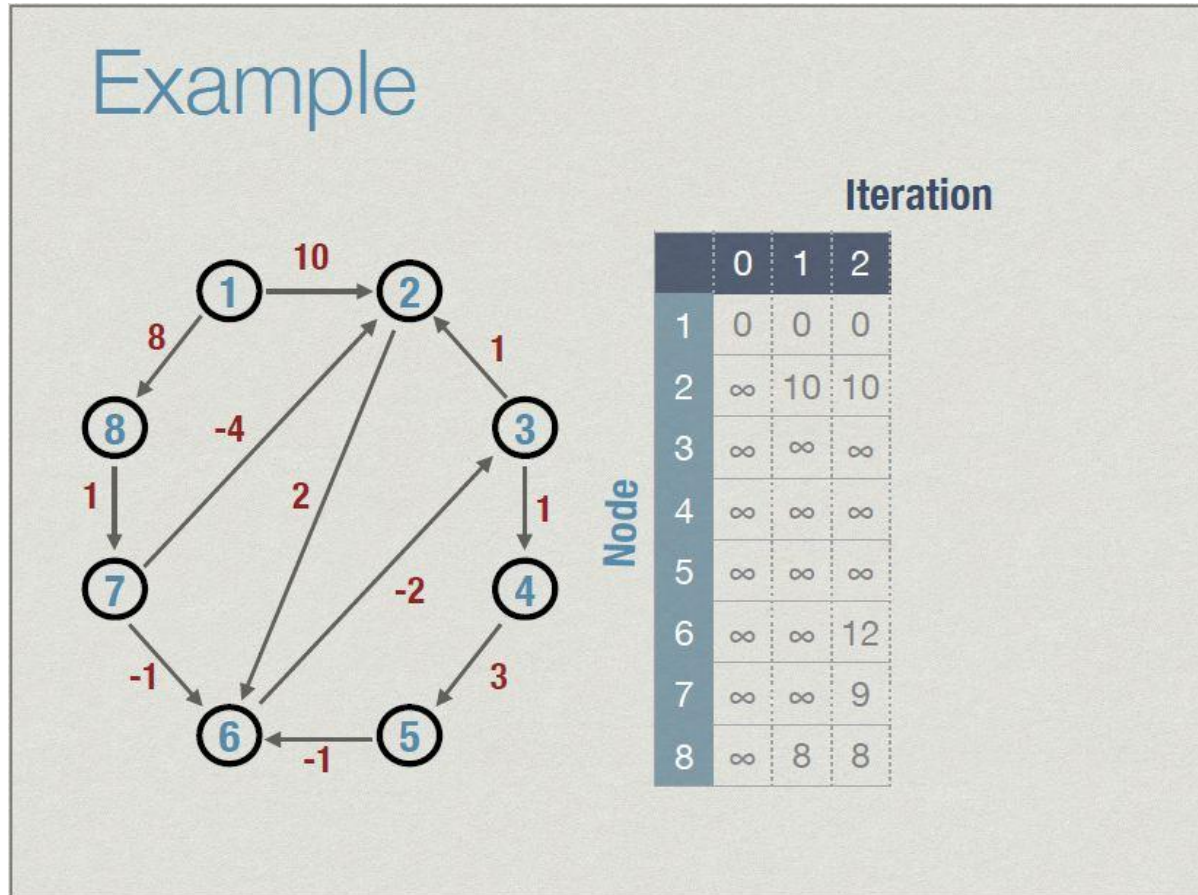
Bellman-Ford algorithm



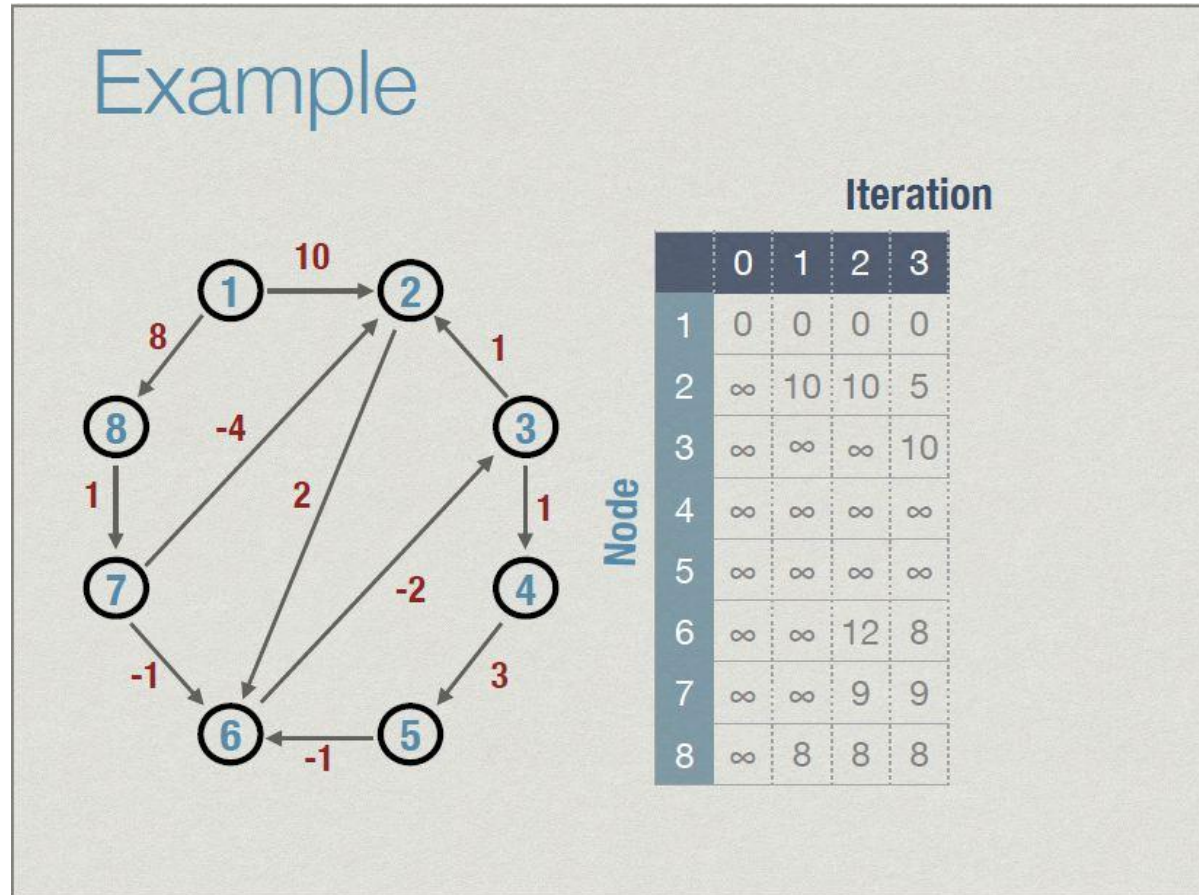
Bellman-Ford algorithm



Bellman-Ford algorithm

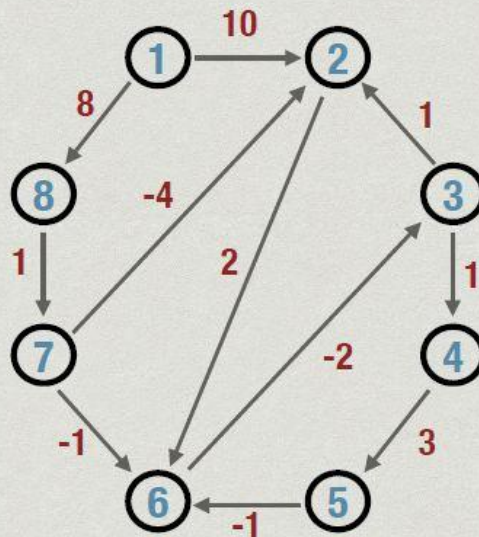


Bellman-Ford algorithm



Bellman-Ford algorithm

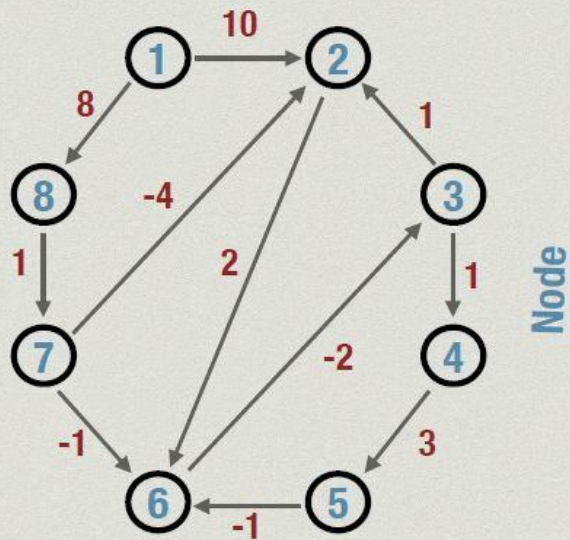
Example



	Iteration				
	0	1	2	3	4
Node 1	0	0	0	0	0
Node 2	∞	10	10	5	5
Node 3	∞	∞	∞	10	6
Node 4	∞	∞	∞	∞	11
Node 5	∞	∞	∞	∞	∞
Node 6	∞	∞	12	8	7
Node 7	∞	∞	9	9	9
Node 8	∞	8	8	8	8

Bellman-Ford algorithm

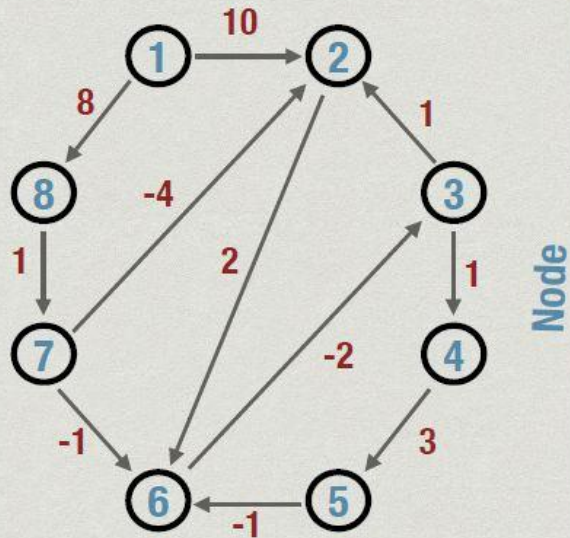
Example



	Iteration					
	0	1	2	3	4	5
Node 1	0	0	0	0	0	0
Node 2	∞	10	10	5	5	5
Node 3	∞	∞	∞	10	6	5
Node 4	∞	∞	∞	∞	11	7
Node 5	∞	∞	∞	∞	∞	14
Node 6	∞	∞	12	8	7	7
Node 7	∞	∞	9	9	9	9
Node 8	∞	8	8	8	8	8

Bellman-Ford algorithm

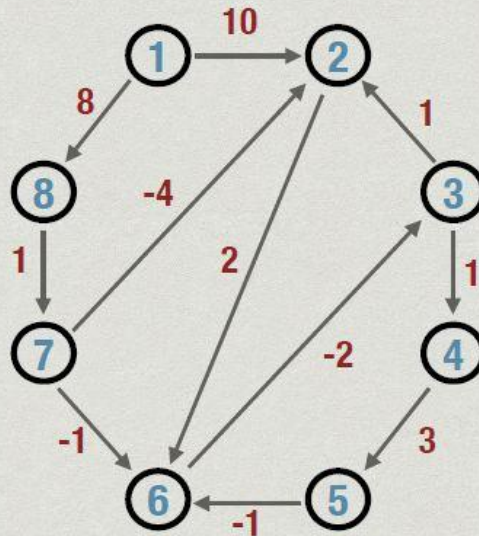
Example



	Iteration						
	0	1	2	3	4	5	6
1	0	0	0	0	0	0	0
2	∞	10	10	5	5	5	5
3	∞	∞	∞	10	6	5	5
4	∞	∞	∞	∞	11	7	6
5	∞	∞	∞	∞	∞	14	10
6	∞	∞	12	8	7	7	7
7	∞	∞	9	9	9	9	9
8	∞	8	8	8	8	8	8

Bellman-Ford algorithm

Example



	Iteration							
	0	1	2	3	4	5	6	7
1	0	0	0	0	0	0	0	0
2	∞	10	10	5	5	5	5	5
3	∞	∞	∞	10	6	5	5	5
4	∞	∞	∞	∞	11	7	6	6
5	∞	∞	∞	∞	∞	14	10	9
6	∞	∞	12	8	7	7	7	7
7	∞	∞	9	9	9	9	9	9
8	∞	8	8	8	8	8	8	8

Bellman-Ford algorithm (Complexity)

- Outer loop runs n times
- In each loop, for each *edge* (j, k) , we run $update(j, k)$
 - Adjacency matrix- $O(n^2)$ to identify all edges
 - Adjacency list - $O(m)$
- Overall
 - Adjacency matrix- $O(n^3)$
 - Adjacency list - $O(mn)$

Weighted graphs

- Negative weights are allowed, but not negative cycles
- Shortest paths are still well defined
- Bellman-Ford algorithm computes single-source shortest paths
- Can we compute shortest paths between all pairs of vertices?

About shortest paths

- Shortest paths will never loop
 - Never visit the same vertex twice
 - At most length $n - 1$
- Use this to inductively explore all possible shortest paths efficiently

Inductively exploring shortest paths

- Simplest shortest path from i to j is a direct edge(i, j)
- General case:
 - $i \rightarrow v_1 \rightarrow v_2 \rightarrow v_3 \dots \rightarrow v_m \rightarrow j$
 - All of $\{v_1, v_2, v_3 \dots, v_m\}$ are distinct, and different from i and j
 - Restrict what vertices can appear in this set

Inductively exploring shortest paths ...

- Recall that $V = \{1, 2, \dots, n\}$
- $W^k(i, j)$: weight of shortest path from i to j among paths that only go via $\{1, 2, \dots, k\}$
 - $\{k+1, \dots, n\}$ cannot appear on the path
 - i, j themselves need not be in $\{1, 2, \dots, k\}$
- $W^0(i, j)$: direct edges
 - $\{1, 2, \dots, n\}$ cannot appear between i and j

Inductively exploring shortest paths ...

- From $W^{k-1}(i, j)$ to $W^k(i, j)$
- *Case 1*: Shortest path via $\{1, 2, \dots, k\}$ does not use vertex k
 - $W^k(i, j) = W^{k-1}(i, j)$
- *Case 2*: Shortest path via $\{1, 2, \dots, k\}$ does go via k
 - k can appear only once along this path
 - Break up as paths i to k and k to j , each via $\{1, 2, \dots, k-1\}$
 - $W^k(i, j) = W^{k-1}(i, k) + W^{k-1}(k, j)$
- Conclusion: $W^k(i, j) = \min(W^k(i, j), W^{k-1}(i, k) + W^{k-1}(k, j))$

Floyd-Warshall algorithm

- W^0 is adjacency matrix with edge weights
 - $W^0[i][j]$ = weight (i, j) if there is an edge (i, j) ,
= ∞ , otherwise
- For k in $1, 2, \dots, n$
 - Compute $W^k(i, j)$ from $W^{k-1}(i, j)$ using
$$W^k(i, j) = \min(W^{k-1}(i, j), W^{k-1}(i, k) + W^{k-1}(k, j))$$
- W^n contains weights of shortest paths for all pairs

Floyd-Warshall algorithm ...

function FloydWarshall

for $i = 1$ to n

 for $j = 1$ to n

$W[i][j][0] = \text{infinity}$

for each edge (i, j) in E

$W[i][j][0] = \text{weight}(i, j)$

for $k = 1$ to n

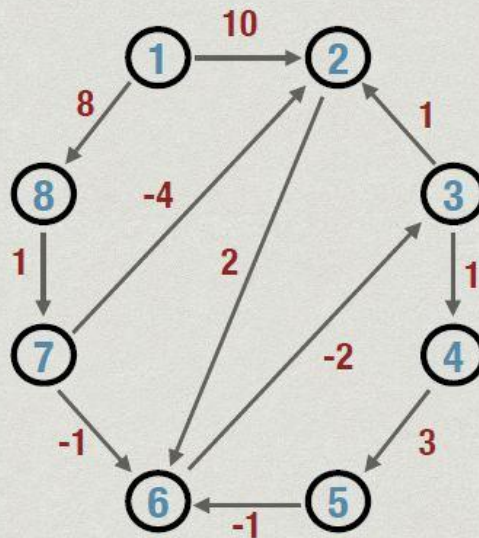
 for $i = 1$ to n

 for $j = 1$ to n

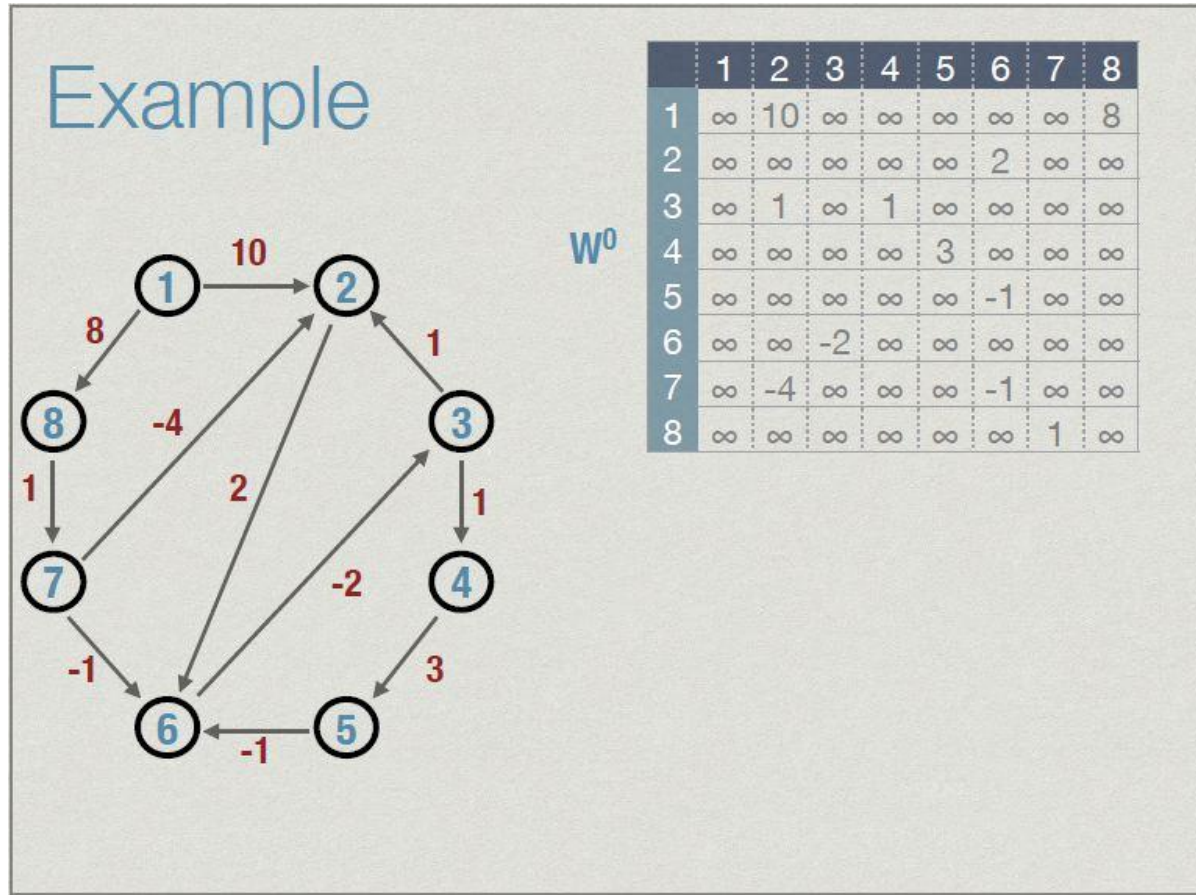
1)) $W[i][j][k] = \min(W[i][j][k - 1], W[i][k][k - 1] + W[k][j][k - 1])$

Floyd-Warshall algorithm ...

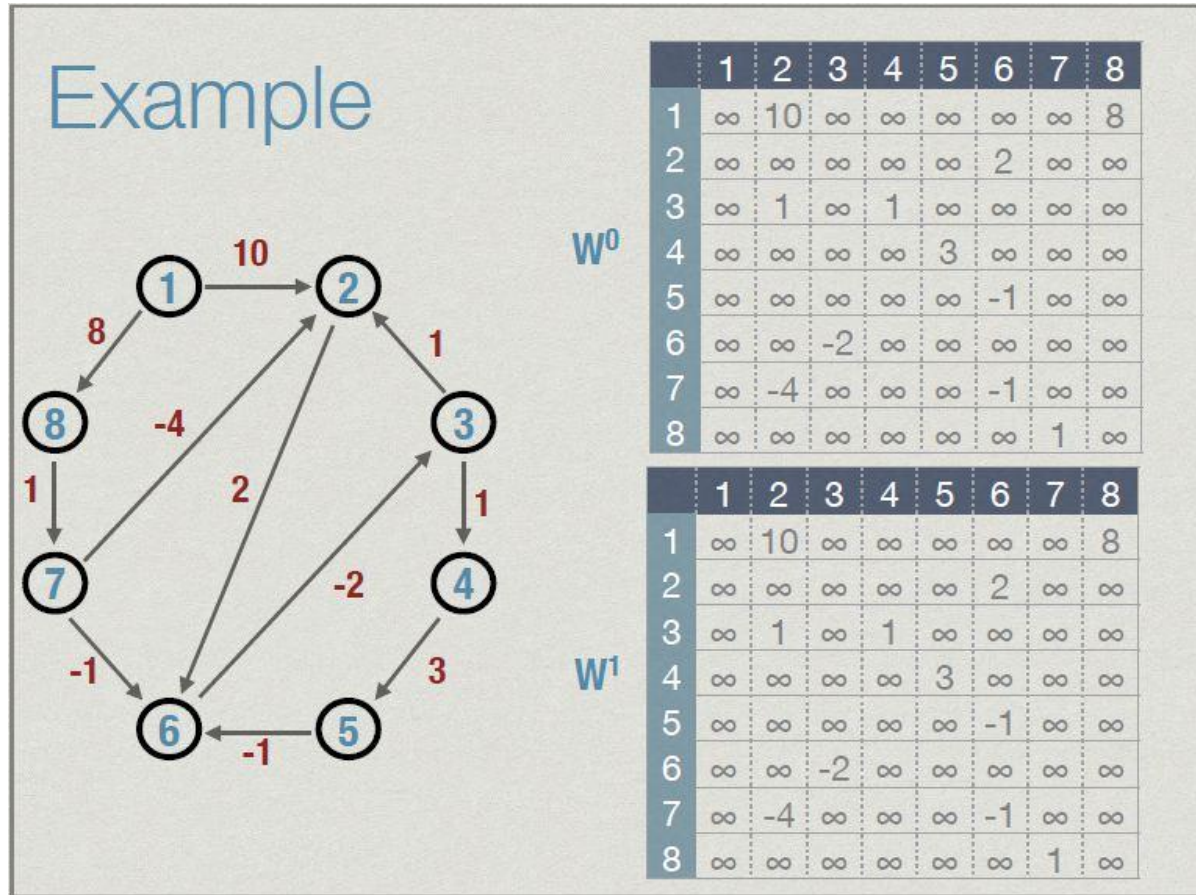
Example



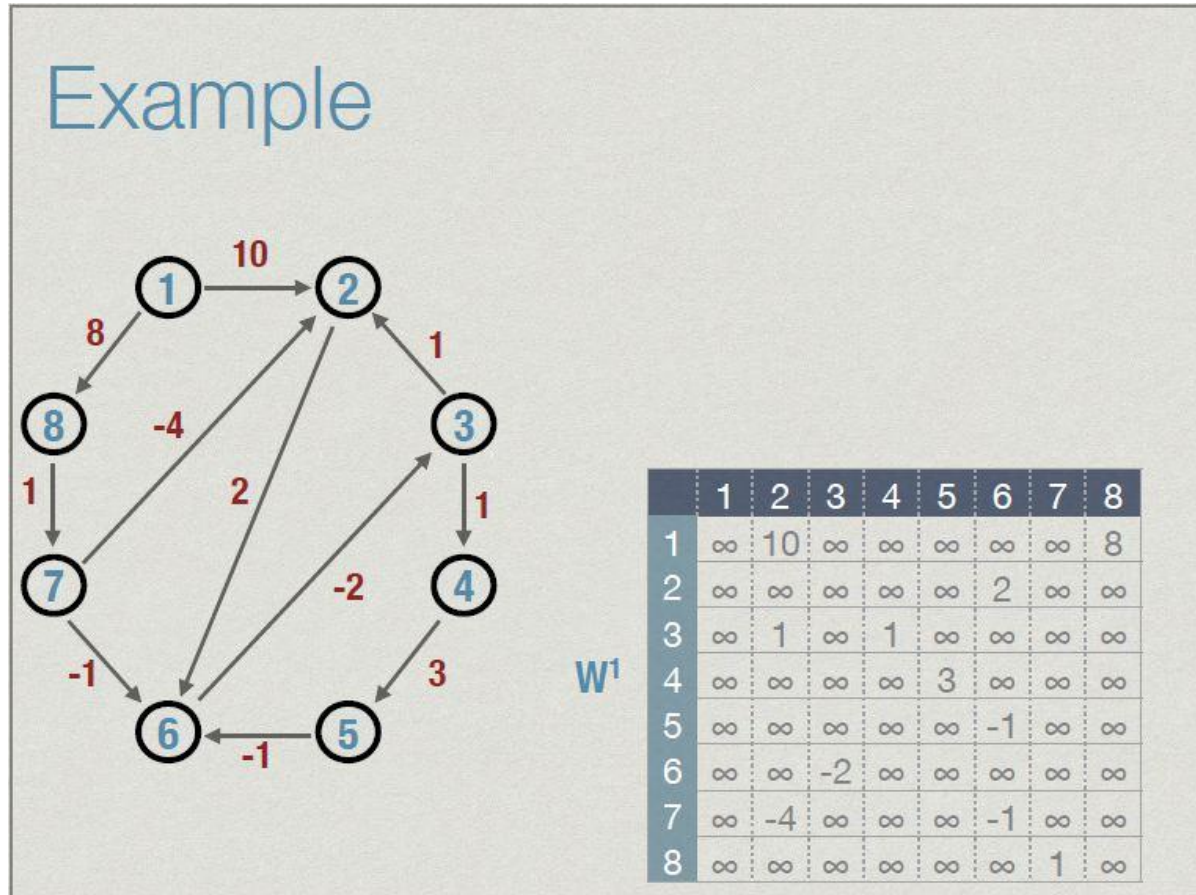
Floyd-Warshall algorithm ...



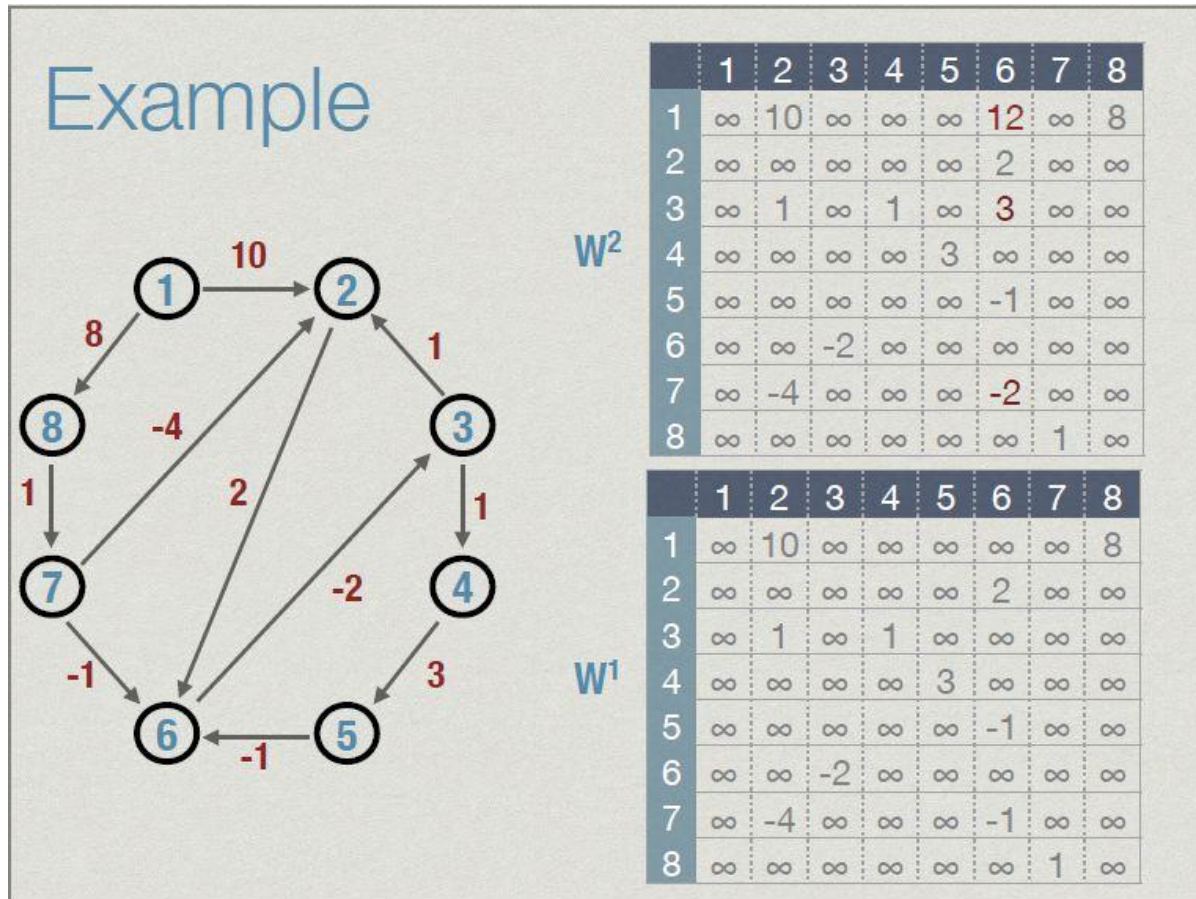
Floyd-Warshall algorithm ...



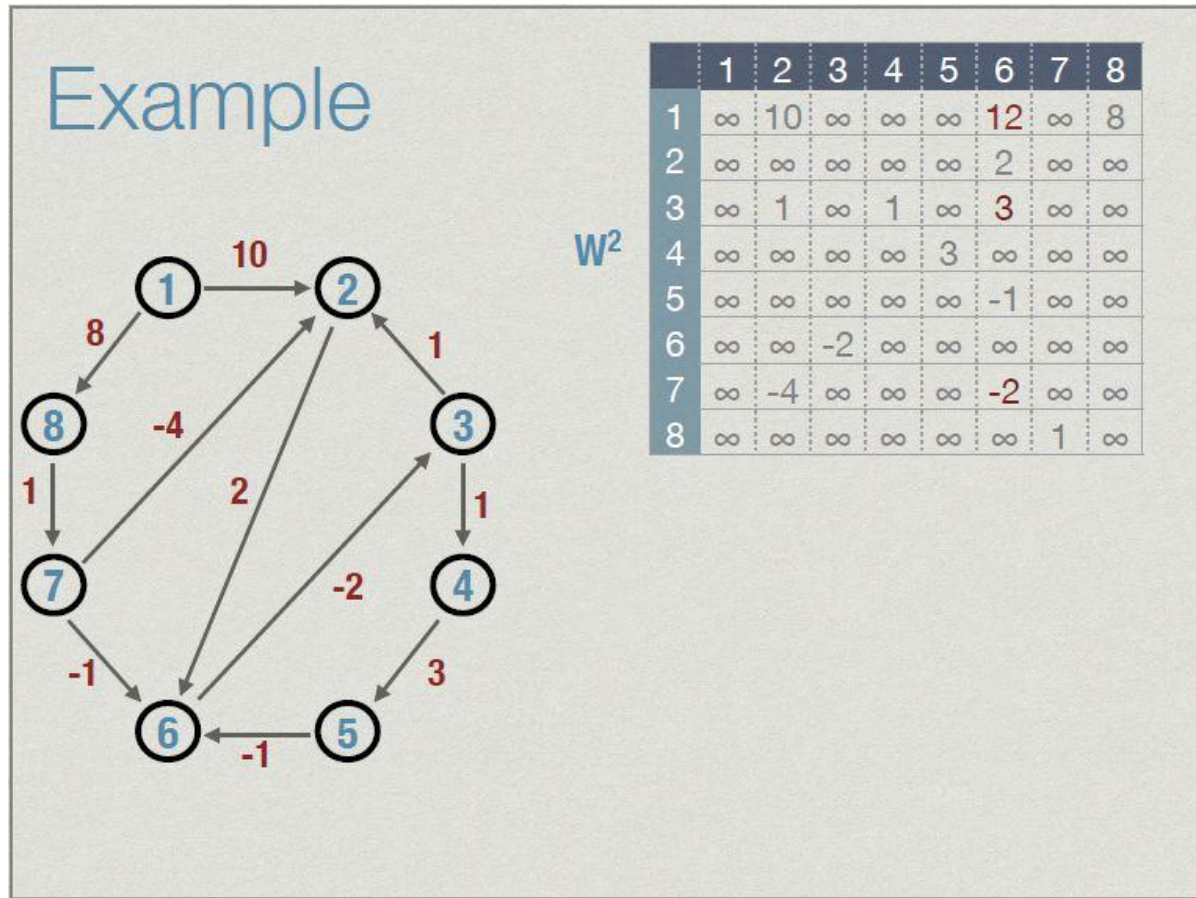
Floyd-Warshall algorithm ...



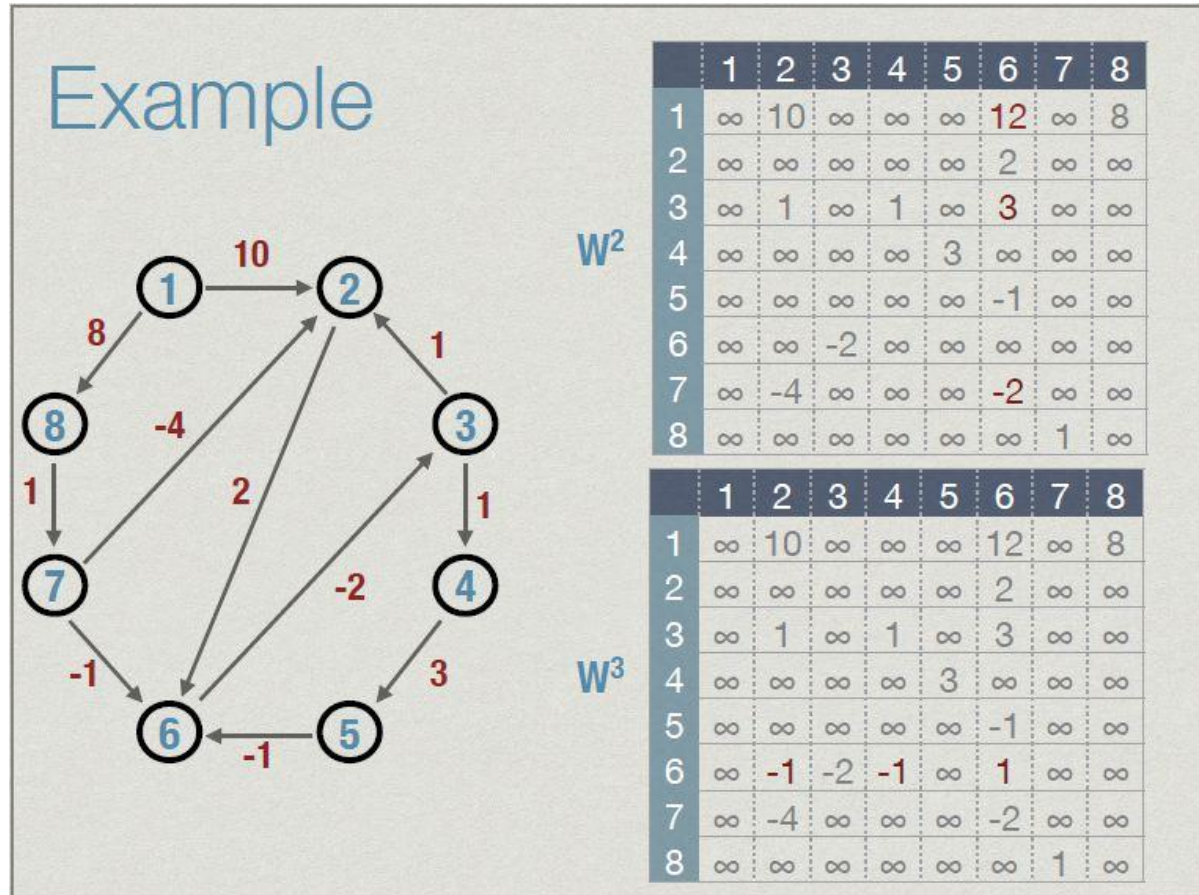
Floyd-Warshall algorithm ...



Floyd-Warshall algorithm ...



Floyd-Warshall algorithm ...



Complexity

Complexity

- Easy to see that the complexity is $O(n^3)$
 - n iterations
 - In each iteration, we update n^2 entries