

Model Checking Using SPIN/PROMELA

Formal Verification :

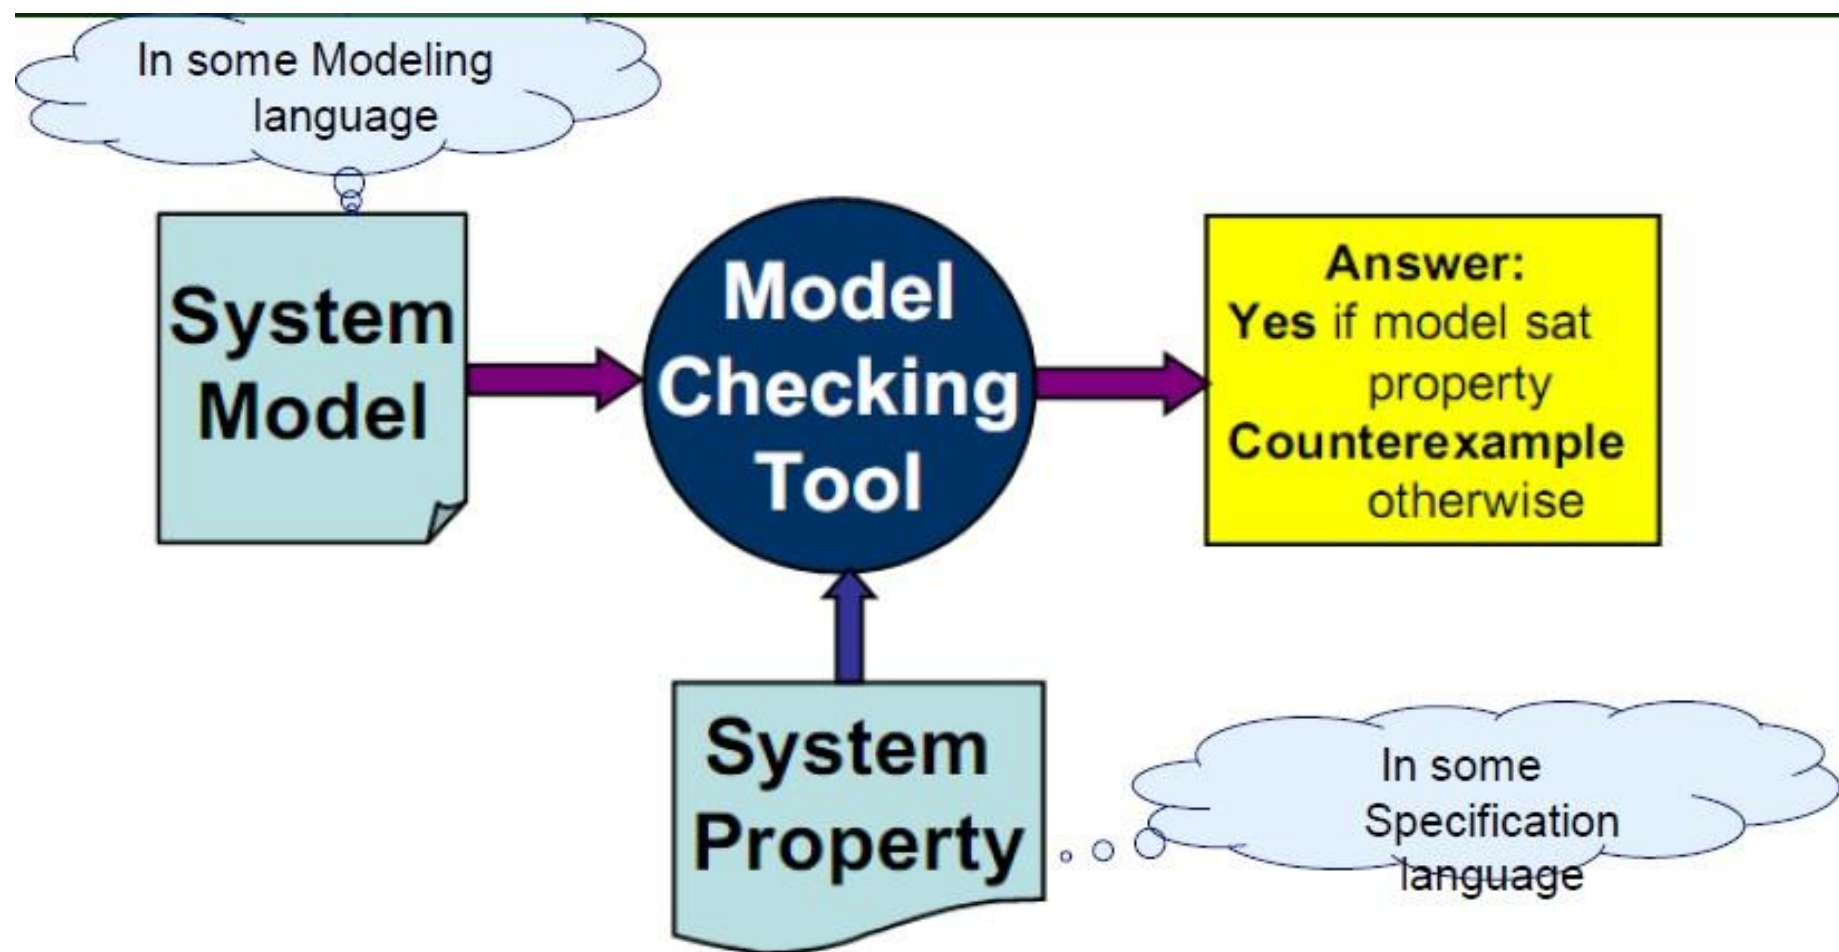
- The process of Formal Verification consists of
 - Requirements
 - Capture Modeling
 - Specification
 - Analysis
 - Documentation
- In practice, some phases may overlap
 - The overall process is iterative rather than sequential

Model Checking

- Model checking
 - Is one of the powerful FORMAL VERIFICATION technique
 - Allows one to verify temporal properties of a finite state representation
 - The finite state representation is that of a typical concurrent system
 - The representation is a model of the system
- The basic idea is
 - That a finite state model of a system is systematically explored in order to determine whether or not a given temporal property holds
 - Deliver a counter-example if the specified property does not hold.

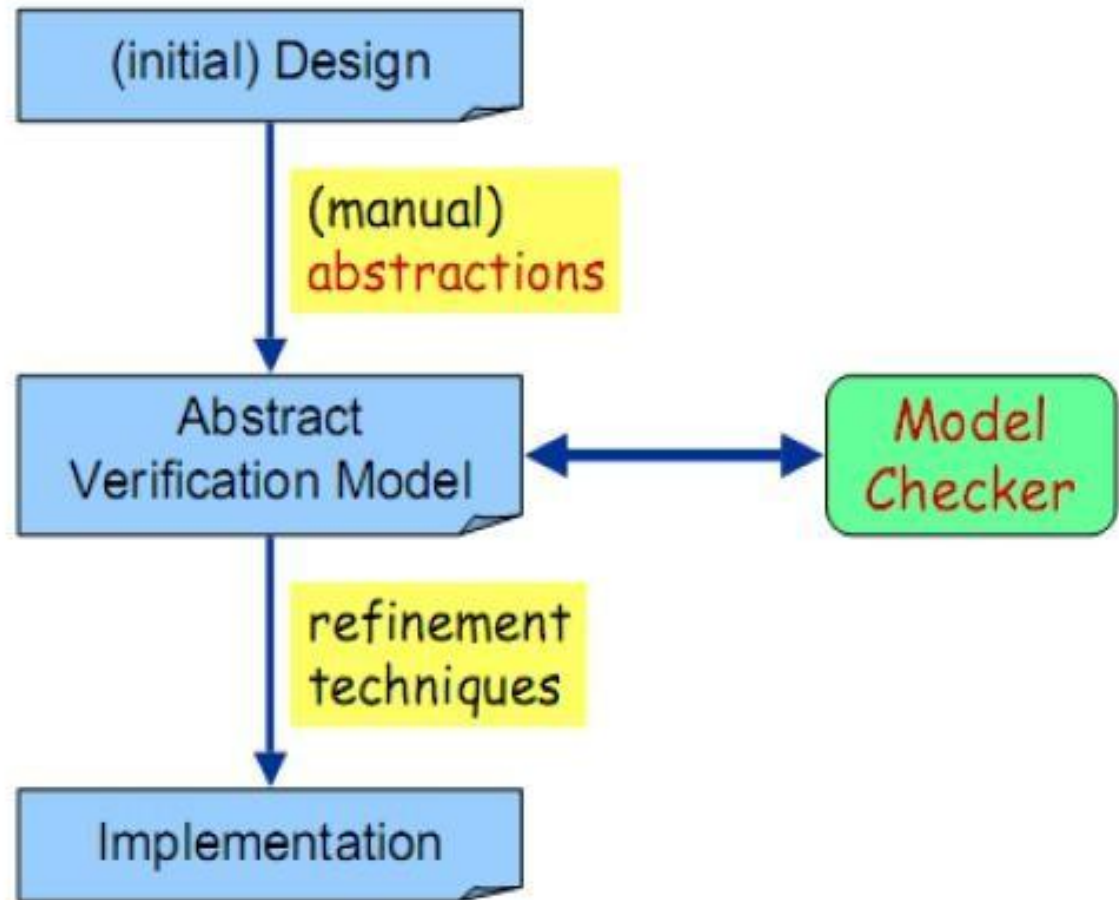
Model Checking

- State space explosion
 - Because a complete state space is generated for a given model, the analysis may fail due to lack of memory, if the model is too large.
 - Can be tackled via abstraction
- Verifying design model
 - Against specification for finite state concurrent systems
- It is an automated technique wherein
 - Inputs
 - Finite state model of the system and properties stated in some standard formalism
 - Outputs
 - Property valid against the model or not



The entire process is automated

Our Focus
here.....



Applicability: Distributed

- Specific concern on the distributed systems
 - Network Applications
 - Data Communication
 - Protocols Multithreaded code
 - Client-Server applications
- Suffer from common design flaws

Common Design Flaws..

- Deadlock
- Livelock,
- Starvation
- Underspecification
 - Unexpected reception of messages
- Overspecification
- Dead code
- Violations of constraints
 - Buffer overruns
 - Array bound violations

Model Checking Definition

- “Model checking is an automated technique that,
 - Given a finite-state model of a system and
 - A logical property, systematically checks whether this property holds for (a given initial state in) that model”

$$M, s \models p$$

- Does system model **M** with initial state **s** satisfy system property **p**
- **M** given as a state machine, that is finite-state
- **p** is usually specified in temporal logic

Model Checking with SPIN

- Involves three
- steps Modeling
 - Convert the design into a formalism to be accepted by the model checking tool SPIN
- Specification
 - State the properties that the design must satisfy
 - Must be complete
- Verification
 - Normally based on a tool i.e. on SPIN
 - Is automatic
 - Analysis of verification results is, however, manual

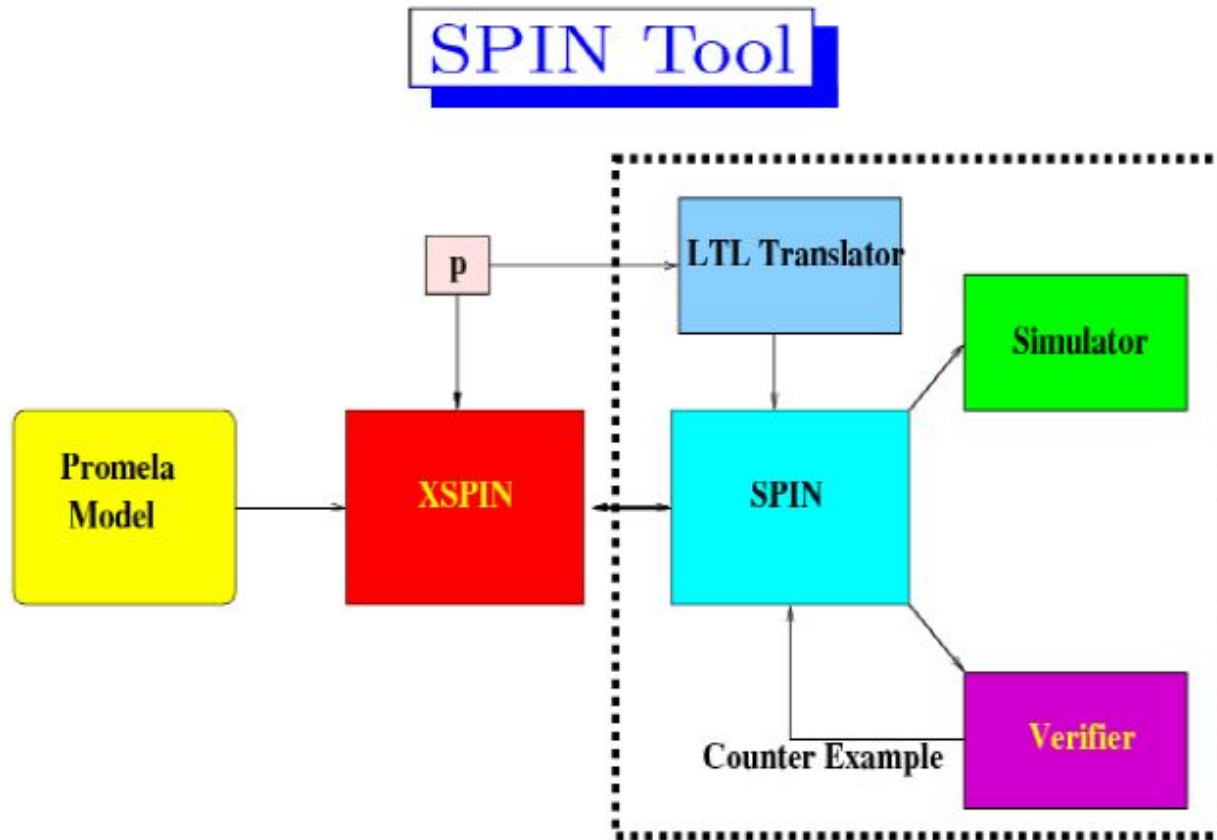
Introduction to

- SPIN – Simple Promela INterpreter
 - A tool for analyzing the logical consistency of concurrent systems, specifically of data communication protocols
 - System model of a concurrent system is described in PROMELA
- PROMELA – PROcess MEta Language
 - Specification language to model finite-state systems & allows dynamic creation of concurrent processes
 - Expressive enough to describe processes and their interactions in Synchronous as well as Asynchronous manner
 - Resembles the programming language C

Introduction to

- PROMELA and SPIN/XSPIN are
 - Developed by Gerard Holzmann at Bell Labs
 - Freeware for non-commercial use
 - Is a State-of-art model checker

Introduction to



Introduction to SPIN

- Major versions:

1.0	Jan 1991	initial version [Holzmann 1991]
2.0	Jan 1995	partial order reduction
3.0	Apr 1997	minimised automaton representation
4.0	late 2002	Ax: automata extraction from C code

- Some success factors of SPIN

- “Press on the button” verification (model checker)
- Very efficient implementation (using C)
- Nice Graphical user interface (Xspin)
- Not just a research tool, but well supported
- Contains more than two decades research on advanced computer aided verification (Many optimization algorithms)

Introduction to

- SPIN's starting page:
 - <http://www.spinroot.com>
 - Basic SPIN manual
 - Getting started with Xspin
 - Getting started with SPIN
 - Examples and Exercises
 - Concise Promela Reference
- Proceedings of all SPIN
 - workshops

Gerard Holzmann's website for papers on
SPIN

PROMELA

- Defining a validation model consisting of
 - A set of states incorporating info about values of variables, program counters etc.
 - A transition relation
- A representation of a FSM in terms of
 - Processes
 - Message Channels
 - State Variables
- Only the design of consistent set of rules to govern interaction amongst processes in a Distributed System NOT the implementation details

PROMELA

- PROMELA model consists of
 - Process declarations
 - Channel declarations
 - Variable declarations
 - Type declarations
 - INIT declarations
 - process

As mentioned, PROMELA model = FSM (Usually a very large)

But it is finite and hence has

- No unbounded data
- No unbounded channels

PROMELA

- A process type (proctype) consists of
 - A name, list of formal parameters, declarations of local variables and body
- A process
 - Executes concurrently with all other processes
 - Communicates with other processes using channels & global variables
 - May access shared variables
 - Defined by proctype declarations
- Each process has
 - Its program counter

PROMELA Variables and Basic Data

- PROMELA variables
 - Provide the means of storing information about the system being modelled
 - May hold global information on the system or information that is local to a particular component
 - Supports five basic data types

Name	Size (bits)	Usage	Range
bit	1	unsigned	0...1
bool	1	unsigned	0...1
byte	8	unsigned	0...255
short	16	signed	$-2^{15} - 1 \dots 2^{15} - 1$
int	32	signed	$-2^{31} - 1 \dots 2^{31} - 1$

PROMELA Variables and Basic Data

- PROMELA variables
 - Variables must be declared before they can be used
 - Variable declarations follow the style of the C programming language
 - A basic data type followed by one or more identifiers and optional initializer
 - byte count, total = 0
 - By default all variables of the basic types are initialized to 0. As in C, 0 (zero) is interpreted as false while any non-zero value is interpreted as true

The init process

- All PROMELA programs must contain an init process
 - Is similar to the main() function within a C program
 - The execution of a PROMELA program begins with the init process
 - An init process
 - Takes the form:
 - `init { /* local declarations and statements */ }`
 - `init {skip}`
 - While a proctype definition declares the behavior of a process, the instantiation and execution of a process definition is coordinated via the init process.

Statement

- A process body consists of sequence of
- statements A statement is either
 - Executable: It can be executed
 - immediately Blocked: It cannot be executed
- An assignment statement is always executable
 - E.g. $x = 2;$

Statement

- An expression is also a statement;
 - It is executable if it evaluates to non-zero
 - skip, $2 < 3$ are always executable
 - $X < 27$ is executable only if the value of x is less than 27
 - A run statement is executable
 - Only if the process can be created
 - Returns 0 if this cannot be done
 - Value otherwise returned is a run-time process ID number
 - `run()` is defined as an operator and so can be embedded in other expressions.

```
proctype A(byte state; short set)
{ (state==1) → state = set
}
init {run A(1,3) }
```

Executability of

- Promela

- Does not make a distinction between a condition and a statement
 - E.g. the simple boolean condition `a == b` represents a statement in Promela
- Statements are either executable or blocked.
 - The execution of a statement is conditional or it being blocked.
- Notion of statement executability provides the basic means by which process synchronization can be achieved.
- E.g.
 - `while (a != b) skip /*Conventional busy wait */`
 - `(a == b) /* Promela equivalent */`

Hello

- A simple two process system
 - `proctype hello() { printf("Hello"); }`
 - `proctype world() { printf ("World \n"); }`
 - `init { run hello(); run world(); }`
- `init` is the starting point
- `run` operator is executable only if process instantiation is possible
- If a `run` is executable then a `pid` is returned. The `pid` for a process can be accessed via the predefined local variable `_pid`.
- The execution of `run` does not wait for the associated process to terminate. i.e. further applications of `run` will be executed concurrently

Process

- A process can be instantiated also by using “active” in front of proctype definition.

- i.e. HelloWorld can also be instantiated as
 - active proctype hello() {printf(“Hello”)}
 - active proctype world()
{printf(“World\n”)}

Multiple instances of the same proctype declaration can be generated using an optional array suffix , e.g.

```
active [4] proctype hello() {printf(“Hello”)}  
active [7] proctype world()  
{printf(“World\n”)}
```

```
proctype Foo(byte x) {  
    ...  
}
```

```
init {  
    int pid2 = run Foo(2);  
    run Foo(27);  
}
```

```
active [3] proctype Bar() {  
    ...  
}
```

- a process can be created at any point in the execution (even within any process)
- processes can also be created using active in front of proctype declaration

Other Data

- Arrays

- An array type is declared as `int table[max]`
- This generates an array of integers i.e. `table[0]`, `table[1]`,...

- Enumerated Types

- A set of symbolic constants is declared as
`mtype = {LINE_CLEAR, TRAIN_ON_LINE, LINE_BLOCKED}`
- A program can only contain one `mtype` declaration which must be global

- Structures

- A record data type is declared as
`typedef msg {byte data[4], byte checksum}`
- Structure access is as in C
`msg message;`
`... message.data[0]`

PROMELA – An illustration

/* a Hello World PROMELA model for SPIN
*/

```
/* A "Hello World" Promela model for SPIN. */  
active proctype Hello() {  
    printf("Hello process, my pid is: %d\n", _pid);  
}  
init {  
    int lastpid;  
    printf("init process, my pid is: %d\n", _pid);  
    lastpid = run Hello();  
    printf("last pid was: %d\n", lastpid);  
}
```

- How many processes are created, here?

Statement


- Two types of statement delimiters
 - ; and ->
 - Use the one that is most appropriate at the given situation Usually ; is used between ordinary statements
 - -> is often used after “guards” in a if OR do statement, pointing at what comes next
 - These can be used interchangeably.

Interleaving

- PROMELA processes execute concurrently
- Non-deterministic scheduling of the processes Processes are interleaved
 - Statements of different processes do not occur at the same time
 - Exception: rendezvous communication
- All statements are atomic;
 - each statement is executed without interleaving with other processes.
- Each process may have several different possible actions enabled at each point of execution
 - Only one choice is made, non-deterministically i.e. randomly

```
init {run A() ; run B();}
```

```
byte x = 2, y = 3;  
proctype A() {x = x + 1}  
proctype B() {  
    x = x - 1;  
    y = y + x  
}
```



Deterministic V/s Nondeterministic

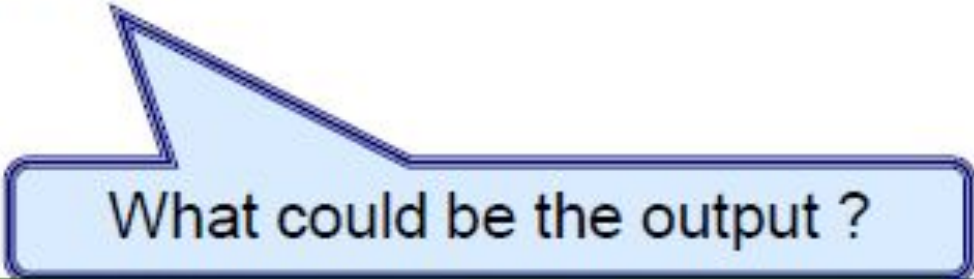
Behaviour

- Deterministic behaviour
 - A process is deterministic if for a given start state, it behaves in exactly the same way ; if supplied with the same stimuli from its environment.
- Non-deterministic behaviour
 - A process is non-deterministic if it need not always behave in exactly the same way ; each time it executes from a given start state with the same stimuli from its environment
 - Hence, race conditions can occur..

Race conditions

- Solutions
 - Use standard mutex
 - algorithms Use atomic sequences

```
byte state = 1;  
proctype A() {      { (state==1) → state=state + 1}}  
proctype B() {      { (state==1) → state=state - 1}}  
init { run A(); run B() }
```

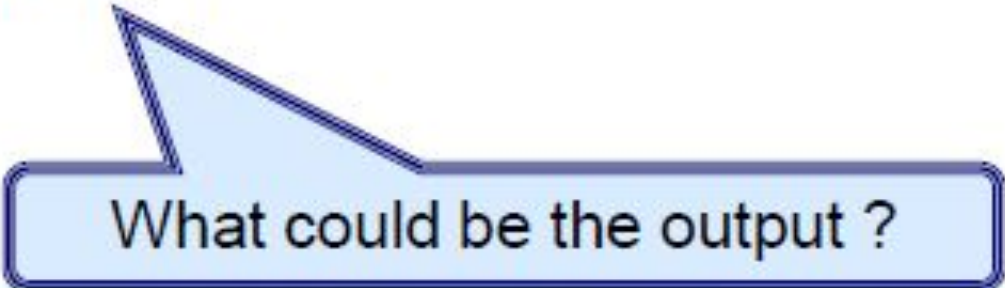


What could be the output ?

Race conditions

- Solutions
 - Use standard mutex algorithms
 - Use atomic sequences

```
byte state = 1;  
proctype A() {atomic{(state==1) → state=state + 1}}  
proctype B() {atomic{(state==1) → state=state - 1}}  
init { run A(); run B() }
```



What could be the output ?

Using atomic

- atomic keyword
 - Helps avoid the undesirable interleaving of the PROMELA execution sequences...
 - Restricts the level of interleaving and so
 - Reduces complexity when it comes to validating a PROMELA model.
 - However, atomic should be used carefully...

PROMELA Control

Structures

- Three ways for achieving control flows
 - Statement sequencing
 - Atomic sequencing
 - Concurrent process execution
- PROMELA supports three additional control flow constructs
 - Case selection
 - Repetition
 - Unconditional Jumps

Case Selection

```
byte count;  
proctype counter()  
{  
    if  
        :: count = count + 1  
        :: count = count - 1  
    fi  
}
```

- Chooses one of the executable choices.
- If no choice is executable, the if-statement is blocked.
 - The executability of the first statement (guard) in each sequence determines whether sequence is executed OR not

Case Selection

- An example of case selection with

```
if
:: (n % 2 != 0) -> n = n + 1;
:: (n % 2 == 0) -> skip;
fi
```

- If there is at least one choice (guard) executable,
 - The if statement is executable and SPIN non-deterministically choose one of the alternatives.
- The operator `->` is equivalent to ;
 - By convention, it is used within if-statements to separate the guards from the statements that follow the guards.

Case Selection

- Guards need not be mutually exclusive

if

:: (x >= y) -> max = x;

:: (y >= x) -> max = y;

fi

- If x and y are equal then
 - The selection of which statement sequence is executed is decided at random, giving rise to non-deterministic choice

Repetition

- An example of repetition involving two statement sequences

do

:: (x >= y) -> x = x - y; q = q + 1;

:: (y > x) -> break;

od

- do statement is similar to the if statement...
 - However, instead of executing a choice once, it keeps repeating the execution.
 - The (always executable) break statement may be used to exit a do-loop statement and transfers control to the end of the loop.

Repetition

- The first statement sequence denotes the body of the loop
 - While the second denotes the termination condition
 - Termination, however, is not always a desirable property of these system, in particular, when dealing with reactive systems

do

```
:: (level > max) -> outlet = open;
```

```
:: (level < min) -> outlet = close;
```

od

```

byte count;
proctype counter()
{
    do
        :: count = count + 1
        :: count = count - 1
        :: (count == 0) -> break
    od
}

```

Name	Size (bits)	Usage	Range
bit	1	unsigned	0...1
bool	1	unsigned	0...1
byte	8	unsigned	0...255
short	16	signed	$-2^{15} - 1 \dots 2^{15} - 1$
int	32	signed	$-2^{31} - 1 \dots 2^{31} - 1$

```
proctype counter()  
{  
  do  
    :: (count != 0) →  
      if  
        :: count = count + 1  
        :: count = count - 1  
      fi  
    :: (count == 0) → break  
  od  
}
```

Unconditional Jump

- PROMELA supports the notion of an unconditional jump via the “goto” statement

```
do
    :: (x >= y) -> x = x - y; q = q + 1;
    :: (y > x) -> goto done;
od;
done:
    skip
```

- “done” denotes a label
 - A label can only appear after a statement
 - A goto, like a skip, is always executable.

Assertions

- An assertion is a statement which can be either true or false
- Interleaving assertion evaluation with code execution provides
 - A simple yet very useful mechanism for checking desirable as well as erroneous behavior with respect to our models
- Assertion: Syntax within PROMELA
 - `assert(<logical-statement>)`
 - E.g. `assert(!(doors == open && lift == moving))`
- Within PROMELA we can express local assertions as well as global system assertions

Global Assertions

- A global assertion is also known as a system invariant
 - Is a property that is true in the initial system state and remains true in all possible execution paths.
- To express a system invariant within PROMELA
 - One must define a monitor process that contains the desired system invariant
- To ensure that the global assertion is checked anypoint during the execution
 - An instance of the monitor process has to be run along with the rest of the system model
- In the case of a simulation the checking is not exhaustive, this is achieved within verification mode

```
bit flag1, flag2;  
byte mutex;
```

```
active proctype A() {  
    flag1 = 1;  
    flag2 == 0;  
    mutex++;  
    mutex--;  
    flag1 = 0;  
}
```

```
active proctype B() {  
    flag2 = 1;  
    flag1 == 0;  
    mutex++;  
    mutex--;  
    flag2 = 0;  
}
```

```
active proctype monitor() { assert mutex != 2};
```

- What could be the eventual value of mutex ?
- Is it really achieved ?
- Is the assertion preserved or violated, here ?

“Invalid End state” in SPIN


```
bit flag1, flag2; byte mutex, turn;
```

```
active proctype A() {  
    flag1 = 1;  
    turn = B_TURN;  
    flag2 == 0 || turn == A_TURN;  
    mutex++;  
    mutex--;  
    flag1 = 0;  
}
```

```
active proctype B() {  
    flag2 = 1;  
    turn = A_TURN;  
    flag1 == 0 || turn == B_TURN;  
    mutex++;  
    mutex--;  
    flag2 = 0;  
}
```

```
active proctype monitor() { assert mutex != 2);
```

First software-only solution to the mutex problem for two processes...

Timeouts

- Reactive systems typically require a means of aborting OR rebooting when a system deadlocks.
- PROMELA provides a primitive statement called timeout for the purpose.

```
proctype watchdog ()  
{ do  
  :: timeout -> guard!reset  
od  
}
```


- The timeout condition becomes true when no other statements within the overall system being modeled are executable

Exceptions

- The unless statement
- A useful exception handling feature
- {statements-1} unless {statements-2}
- Consider an alternate watchdog process.

Message Channels

- Means to achieve communication between distinct processes?
- PROMELA supports message channels:
 - Provide a more natural and sophisticated means of modelling inter-process communication/data transfer
- Channel declaration:
 - `chan <name> = <dim> of {<t1>,<t2>,<t3> ... <tn>}`
- E.g. `chan q = [1] of {byte}` ; `chan q = [3] of {mtype, int}`
- If `dim = 0` then synchronous. E.g. `chan q = [0] of {bit}`
- A channel can be defined to be either local or global



Message Channels... Sending-in

- Sending messages through a channel – FIFO buffer (dim>0)
 - Achieved by ! Operator
 - E.g. in_data ! 4
 - Type of the channel and variable must match.
- If multiple data values are to be transferred via each message
 - out_data ! x+1, true, in_data
- The executability of a send statement is dependent upon the associated channel being non-full

Message Channels... Receiving from

- Receiving messages is achieved by ? Operator
 - E.g. `in_data ? Msg`
- If the channel is not empty, the first message is fetched from the channel and is stored in `msg`
- Multiple values can also be fetched.
 - E.g. `out_data ? value1, value2, value 3;`
- The executability of a receive statement is dependent upon the associated channel being non-empty,

Message Channels...

- If more data values are sent per message than can be stored by a channel then the extra data values are lost

- E.g. `in_data ! msg1, msg2 ;` `msg2` will be lost

If fewer data values are sent per message than are expected, then the missing data values are undefined.

E.g. `out_data ! 4, true` and `out_data? x, y , z`

- `x` and `y` will be assigned the values `4` and `true` respectively while the value of `z` will be undefined.

Message Channels...

- len operator
 - To determine the number of messages in a channel
 - E.g. `len(in_data)`
 - If the channel is empty then the statement will block.
- empty, full operators
 - Determine whether or not messages can be received or sent respectively.
 - E.g. `empty(in_data)` ; `full(in_data)`
- Non-destructive retrieve
 - `out_data ? [x, y, z]`
 - Returns 1 if `out_data ? x,y,z` is executable otherwise 0.
 - No side-effects. Only evaluation, not execution. No message retrieved.

Channels as

- Output?

```
proctype A(chan q1) {
    chan q2;
    q1?q2;
    q2!123
}
proctype B(chan qforb) {
    int x;
    qforb?x;
    printf("x = %d\n", x)
}
init {
    chan qname[2] = [1] of { chan };
    chan qforb = [1] of { int };
    run A(qname[0]); run B(qforb);
    qname[0]!qforb
}
```

Communication type

- What was the type of communication pattern observed in the examples till now?
 - Synchronous
 - OR
 - Asynchronous ?
- Why?

Synchronous

- When the channel declaration is
 - `chan ch = [0] of {bit, byte};` i.e. when `dim = 0`
- If `ch ! x` is enabled and
 - If there is a corresponding receive `ch ? X`
 - that can be executed simultaneously and both the statements are enabled
 - Both statements will handshake and together do the transition
- `chan ch = [0] of {bit, byte}`
- P wants to perform `ch ! 1, 3 +`
- 7 Q wants to perform `ch ? 1, x`
- Then after communication x will be 10.

Alternating Bit

Protocol

- To every message, the sender adds a bit
- The receiver acknowledges each message by sending the received bit back.
- The receiver only expects messages with a bit that is expected to receive
- If the sender is sure that the receiver has correctly received the previous message, it sends a new message and it alternates the accompanying bit